



# Neural Networks on Chip Design from the User Perspective


Yu Wang @ VLSID 2020

Prof. of E.E. Dept., Tsinghua University

Co-founder of DeePhi Tech (now part of Xilinx)


[yu-wang@tsinghua.edu.cn](mailto:yu-wang@tsinghua.edu.cn)

<https://nicsefc.ee.tsinghua.edu.cn>



# Deep Learning for Everything

**Rocket: Deep learning**




**Fuel: Big data**

**Engine: Computing platform**

According to Prof. Andrew Ng


Training

– How accurate the model can be

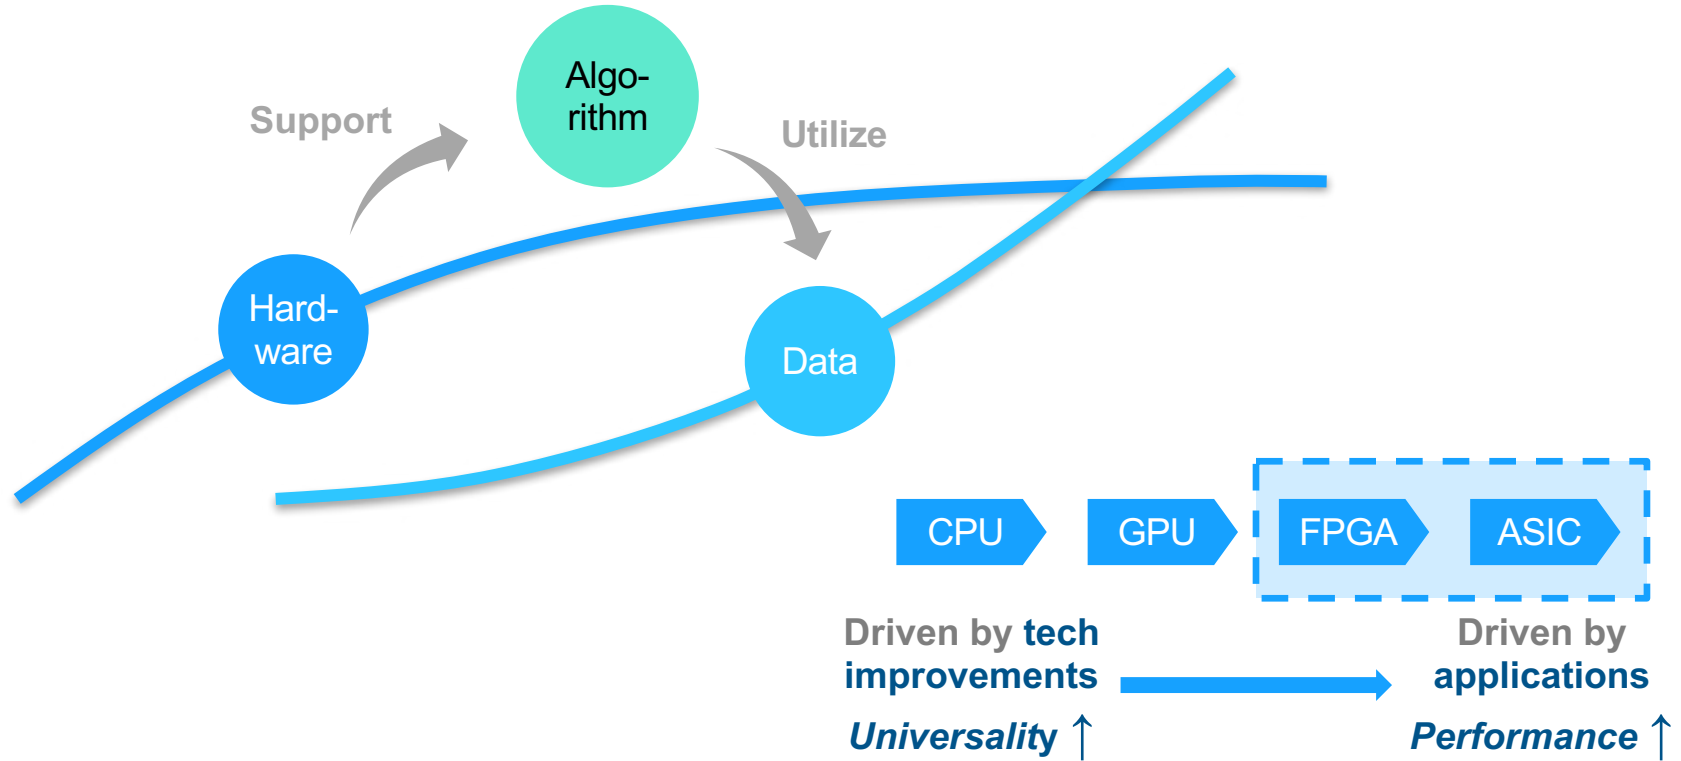


**Inference**

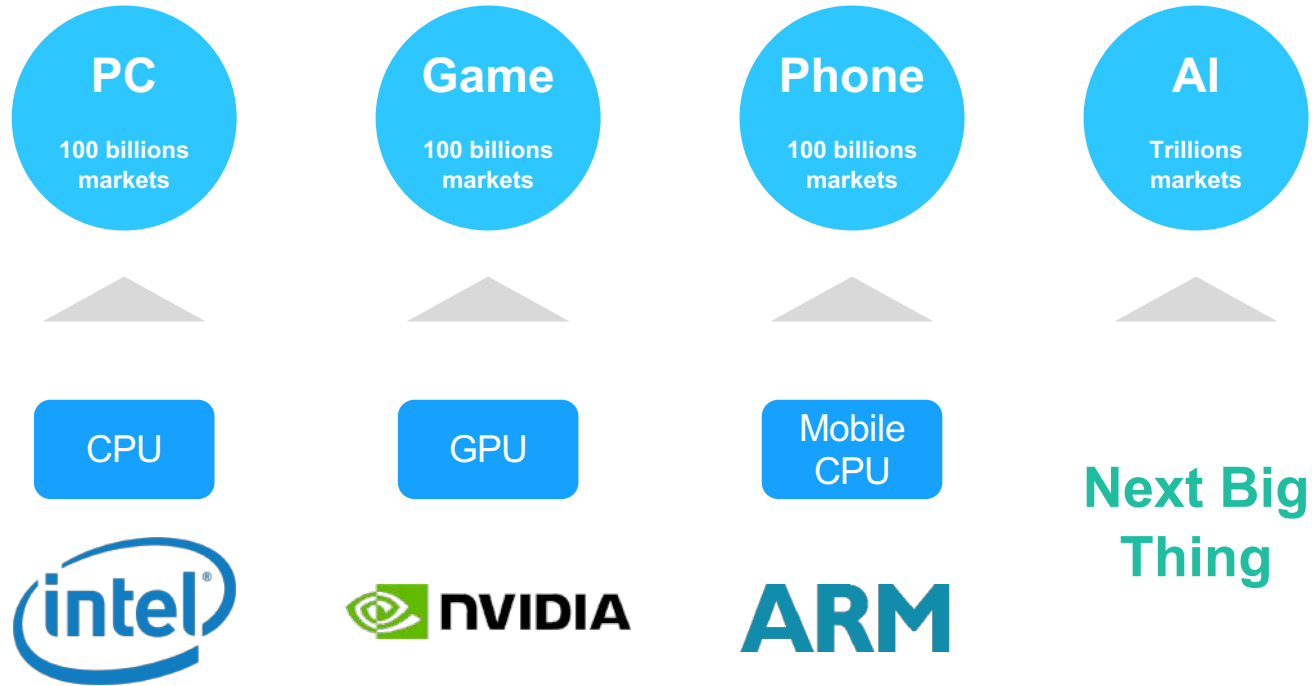
– How many applications can use DL



# Current Hardware is Insufficient to Support Algorithm and Data



# The New Era is Waiting for the Next Rising Star



Back to 2012...

---

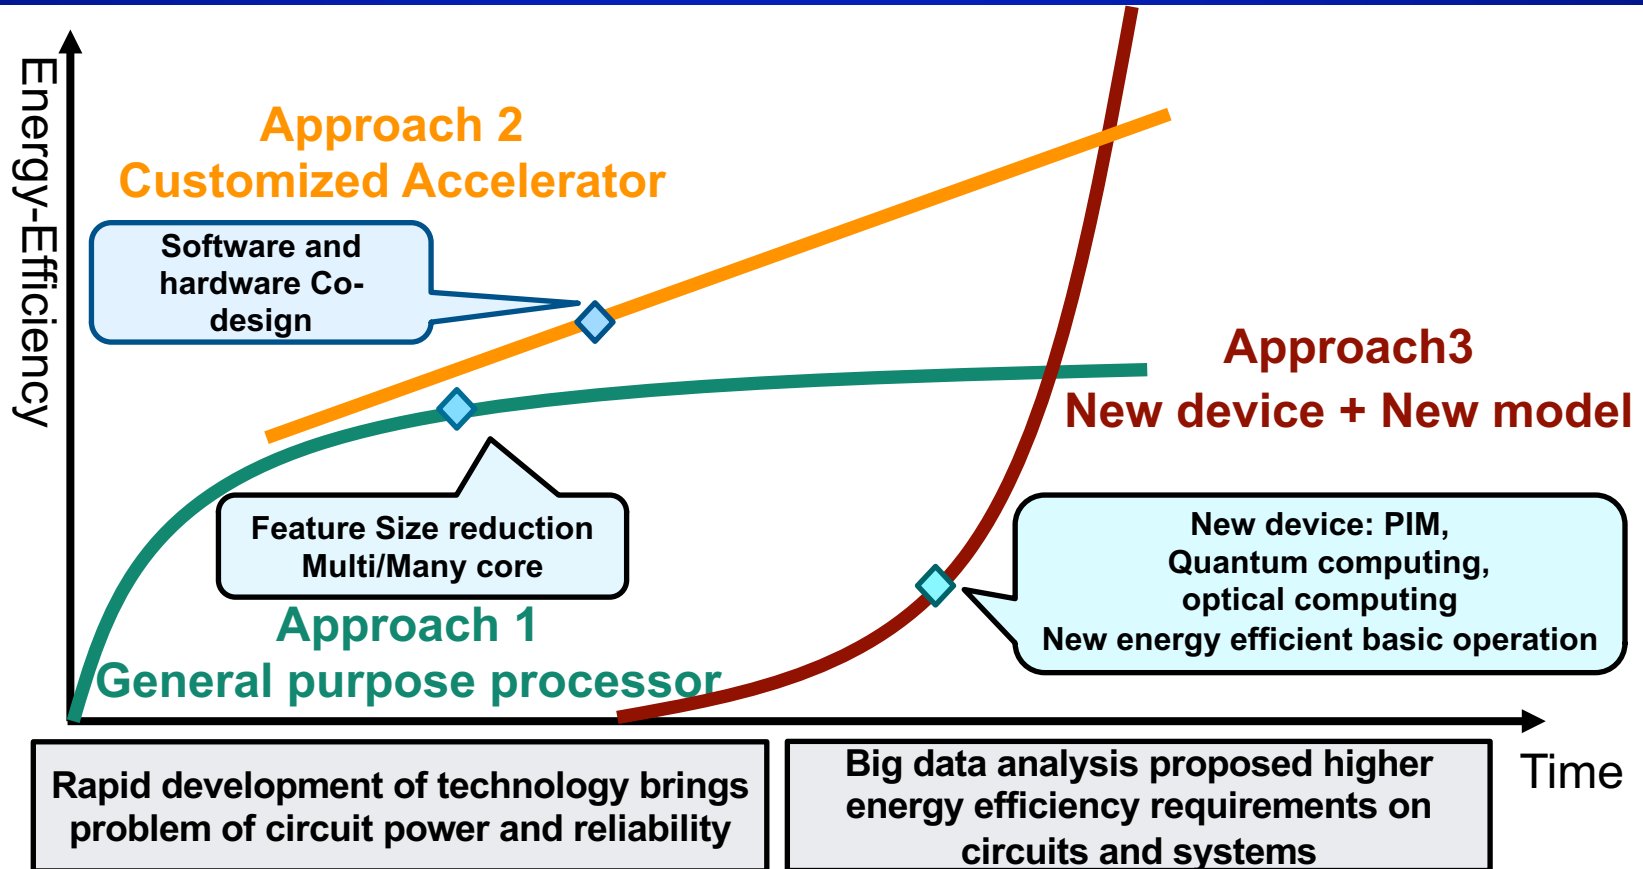
# Are GPUs Enough?

Good, but not good enough....

# Why? Power Consumption and Latency Are Crucial

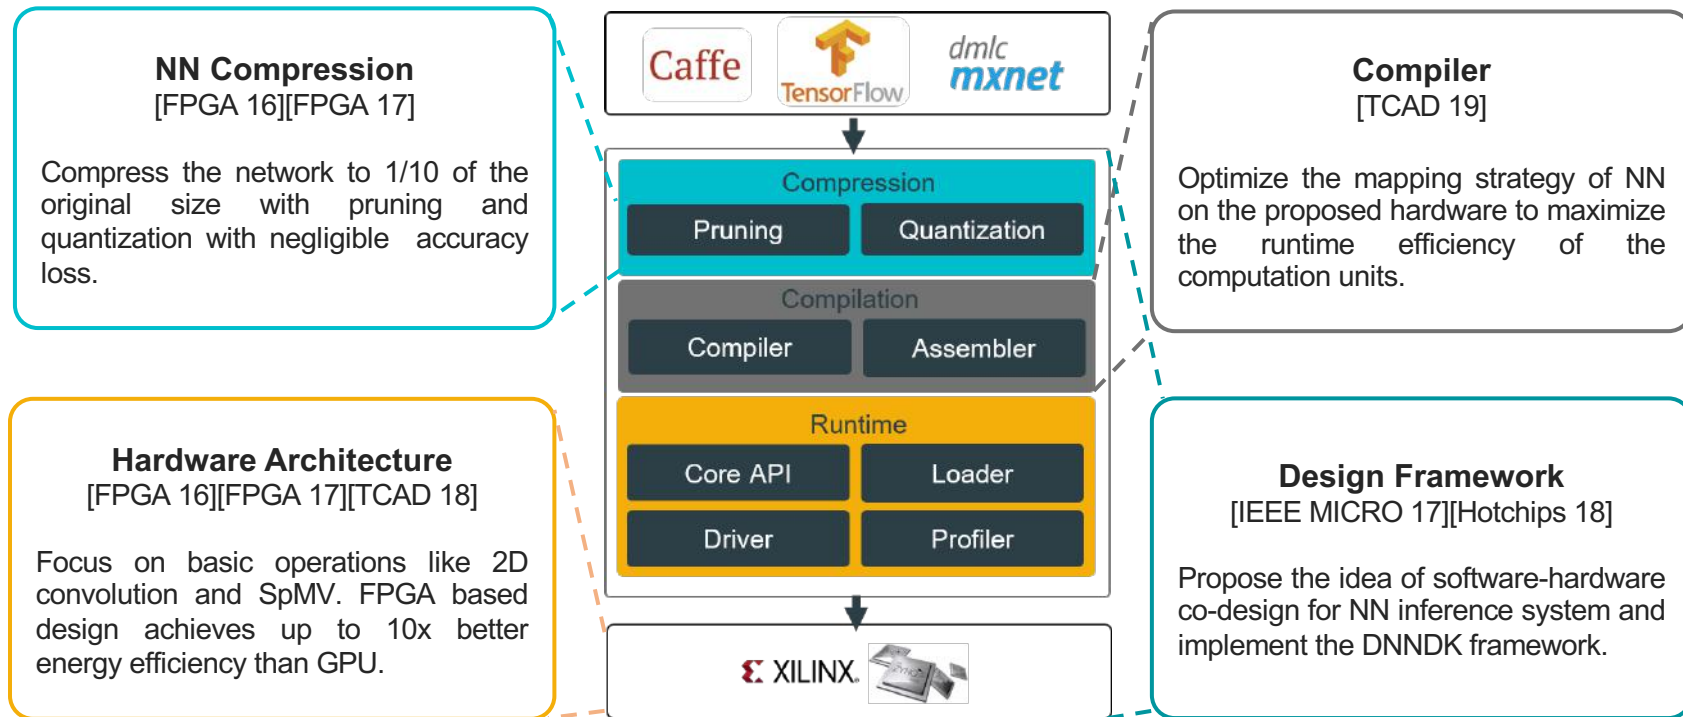


# Development of Energy-Efficient Computing Chips



# Our Previous Work: Software-Hardware Co-design for Energy Efficient NN Inference System

Peak performance ↑ + Utilization ↑ + Workload ↓

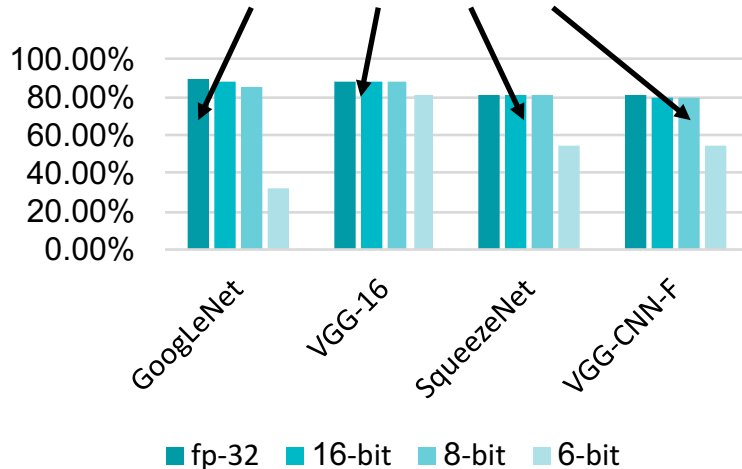




# NN Compression: Quantization

Peak performance ↑ + Utilization ↑ + Workload ↓

Negligible accuracy loss with  
8-bit fixed-point weight & activation



Save FPGA resources

	Xilinx Logic				Xilinx DSP		
	multiplier		adder		multiply & add		
	LUT	FF	LUT	FF	LUT	FF	DSP
fp32	708	858	430	749	800	1284	2
fp16	221	303	211	337	451	686	1
fixed32	1112	1143	32	32	111	64	4
fixed16	289	301	16	16	0	0	1
fixed8	75	80	8	8	0	0	1
fixed4	17	20	4	4	0	0	1


[1] [IEEE MICRO 17] Kaiyuan Guo, Song Han, Song Yao, et al., Software–Hardware Codesign for Efficient Neural Network Acceleration , in IEEE Micro, vol.37, No.2, 2017, pp.18-25.

[2] [IEEE TCAD 18] Kaiyuan Guo, Lingzhi Sui, Jiantao Qiu, Jincheng Yu, Junbin Wang, Song Yao, Song Han, Yu Wang, Huazhong Yang, Angel-Eye: A Complete Design Flow for Mapping CNN onto Embedded FPGA , in IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD), vol.37, No.1, 2018, pp.35-47


# NN Compression: Pruning

Peak performance  $\uparrow$  + Utilization  $\uparrow$  + **Workload**  $\downarrow$

Negligible WER increase  
with 10% weights



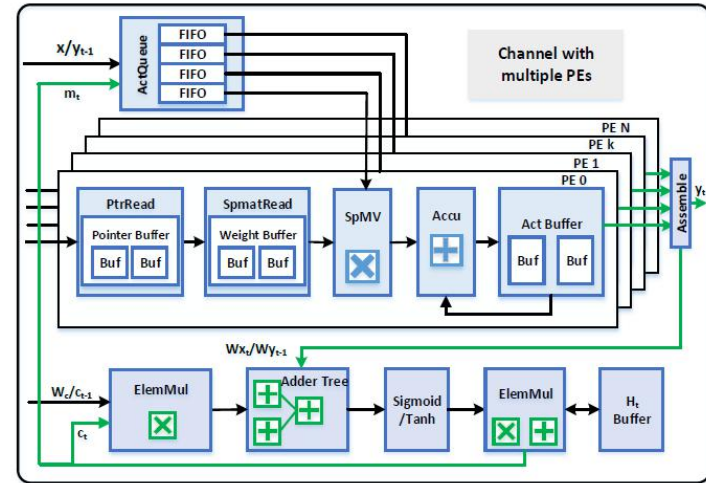
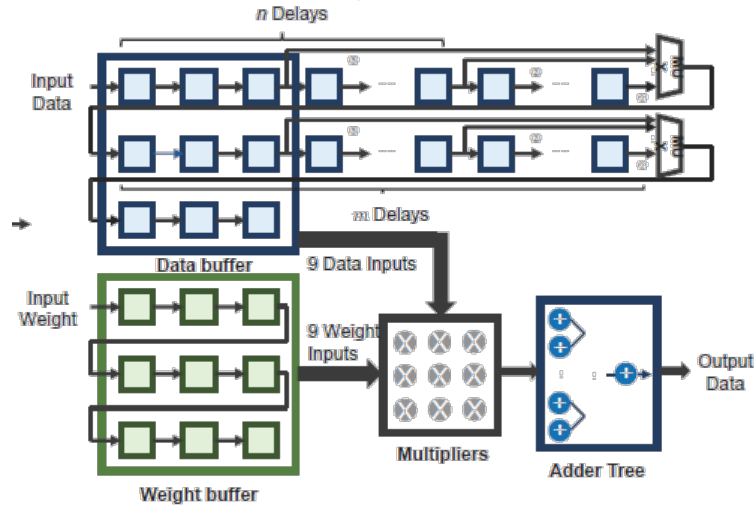
6x acceleration with  
customized hardware



[1] [FPGA 17] Song Han, Junlong Kang, Huizi Mao, et al., ESE: Efficient Speech Recognition Engine with Compressed LSTM on FPGA, in ACM International Symposium on FPGA, 2017, pp.75-84.

# Hardware Architecture – Utilization ↑

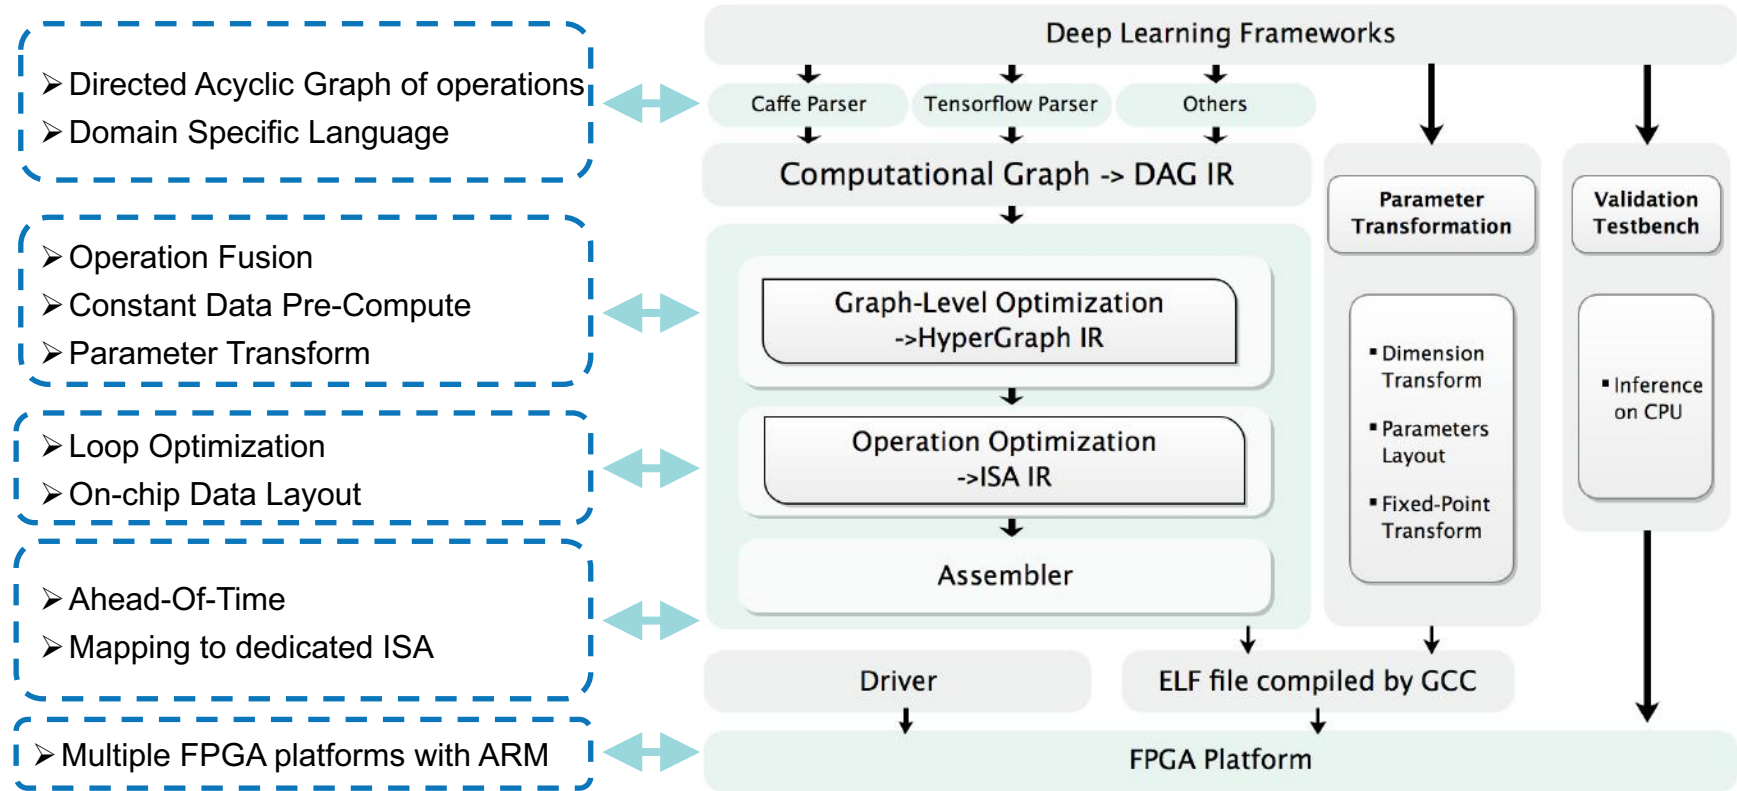
- **Dedicate Datapath:** pipelined MAC array and nonlinear operations, SpMV
- **Memory Architecture:** Line buffer for convolution; scratchpad and pointer buffer for sparse memory access



[1] [FPGA 16] Jiantao Qiu, Jie Wang, Song Yao, et al., Going Deeper with Embedded FPGA Platform for Convolutional Neural Network , in ACM International Symposium on FPGA, 2016, pp.26-35.

[2] [FPGA 17] Song Han, Junlong Kang, Huizi Mao, et al., ESE: Efficient Speech Recognition Engine with Compressed LSTM on FPGA , in ACM International Symposium on FPGA, 2017, pp.75-84.

# Compiler – Utilization ↑



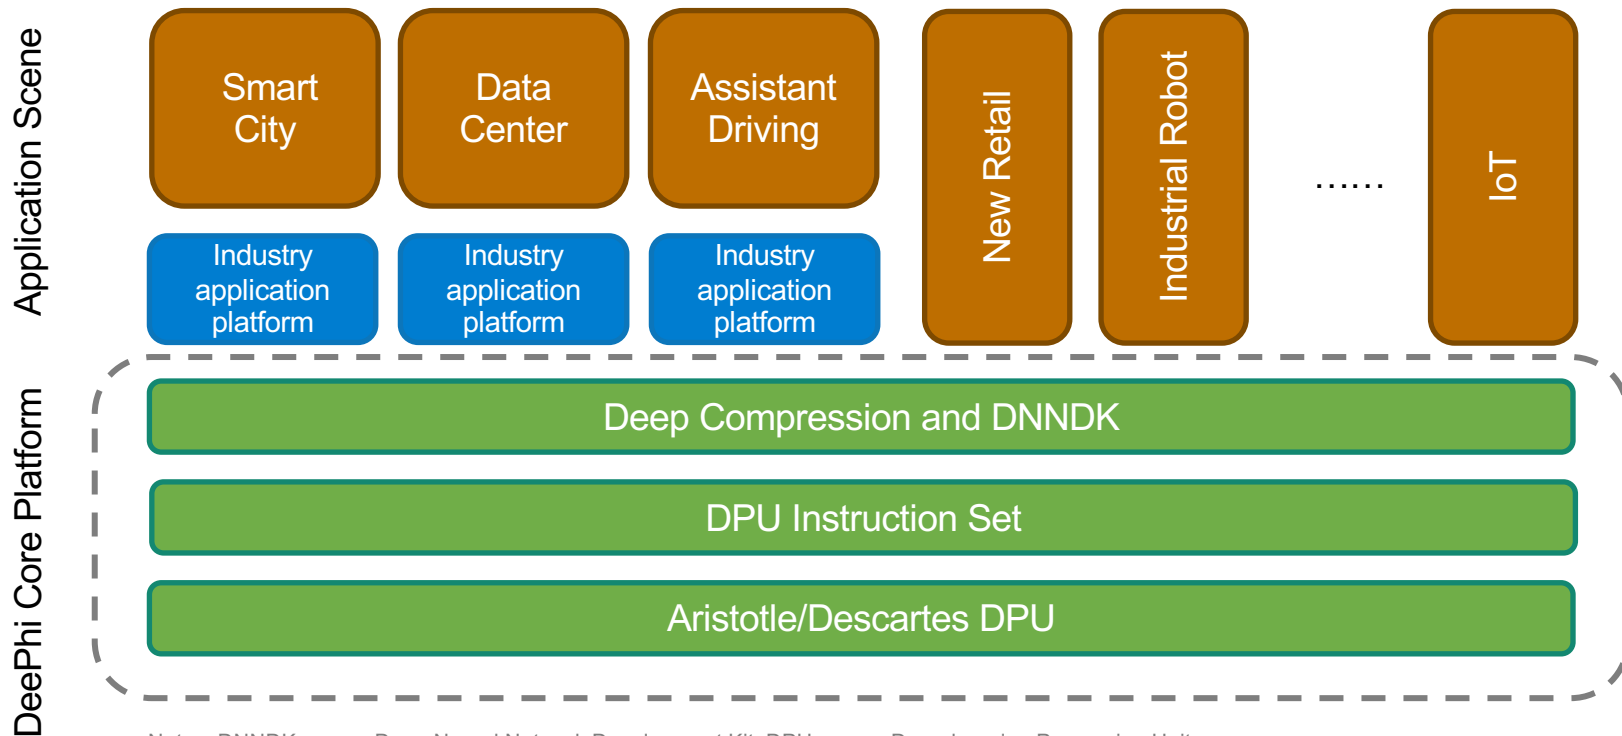
# Compiler

Benchmark	Node Size	Graph Generation(ms)	Isomorphism Fusion(ms)	Compilation Jobs(s)	Evaluation (ms)	Auto Tuning(ms)	Baseline (ms)	After Fusion(ms)	Speedup
Vgg16	32	0.389	9.28	28.266	189.833	0.647	94.312	91.89	2.57%
Resnet50	120	4.112	30.893	29.125	55.097	12.309	39.408	33.631	14.66%
Resnet152	358	18.125	145.023	72.151	141.553	70.53	105.887	92.189	12.94%
GoogleNet	137	3.424	66.069	22.139	28.076	6.61	16.978	11.713	31.01%

- Up to 30% speedup with optimized scheduling
- Low compilation time overhead: < 100s for each network

# DeePhi Tech: From Academy to Industry

- Aims to provide efficient, convenient & economic DL platform for AI applications



Notes: DNNDK means Deep Neural Network Development Kit, DPU means Deep-learning Processing Unit



Now  
Part of



XILINX®

Discovering the Philosophy behind  
Deep Learning Computing



Applications

Products

Developers

Support

About



[Xilinx - Adaptable. Intelligent.](#) > [News](#) > [Press](#) > [2018 Press Releases](#) > [Xilinx Announces the Acquisition of DeePhi Tech](#)

## Xilinx Announces the Acquisition of DeePhi Tech

### Deal to Accelerate Data Center and Intelligent Edge Applications

Jul 17, 2018

**BEIJING and SAN JOSE, Calif., July 17, 2018** – [Xilinx, Inc.](#) (NASDAQ: XLNX) the leader in adaptive and intelligent computing, announced today that it has acquired [DeePhi Technology Co., Ltd](#) (DeePhi Tech), a Beijing-based privately held start-up with industry-leading capabilities in machine learning, specializing in deep compression, pruning, and system-level optimization for neural networks.

DeePhi Tech has been developing its machine learning solutions on Xilinx technology platforms, and the two companies have worked closely together since DeePhi Tech's inception in 2016. DeePhi Tech's neural network pruning technology has been optimized to run on Xilinx FPGAs, enabling breakthrough performance with best-in-class energy efficiency. Xilinx has been a major investor in DeePhi Tech, alongside other prominent international investors, since 2017.

# DPU and DNNDK available now!

## Edge AI Resources

The following resources are available to help you start developing

### Edge AI Tools

Product	Documentation
AI SDK	AI SDK User Guide (UG1354) v2.0
	AI SDK Programming Guide (UG1355) v2.0
DNNDK	DNNDK User Guide (UG1327) v1.6
	DNNDK User Guide (UG1327) v1.5
	DNNDK User Guide (UG1327)_v1.4
DNNDK for SDC	DNNDK User Guide for SDC (UG1331)

### Edge AI Evaluation Boards

Product	Documentation	Image Download
ZCU102 Kit	ZCU102 User Guide (UG1182)	petalinux-user-image-zcu102-zynqr-20190802.img.gz
		xilinx-zcu102-prod-dpu1.4-2018.3-desktop-buster-2019-04-24.img.zip
		2018-12-04-zcu102-desktop-stretch.img.zip
ZCU104 Kit	ZCU104 User Guide (UG1267)	petalinux-user-image-zcu104-zynqmp-sd-20190802.img.gz
		xilinx-zcu104-prod-dpu1.4-desktop-buster-2019-04-23.img.zip
		2018-12-04-zcu104-desktop-stretch.img.zip

### Edge AI Targeted Reference Designs (TRD)

Product	Documentation	Image Download	DNNDK Version
DPU TRD v3.0	DPU Product Guide (PG338 v3.0)	zcu102-dpu-trd-2019-1-190809.zip	v3.1
DPU TRD v2.0	DPU IP Product Guide (PG338 v2.0)	zcu102-dpu-trd-2018-2-190531.zip	v3.0
DPR TRD v1.0	DPU IP Product Guide (PG338 v1.0)	zcu102-dpu-trd-2018-2-190322.zip	v2.08

**DNNDK (Tool) + Ready to use image (DPU on boards) + Targeted Reference Design**

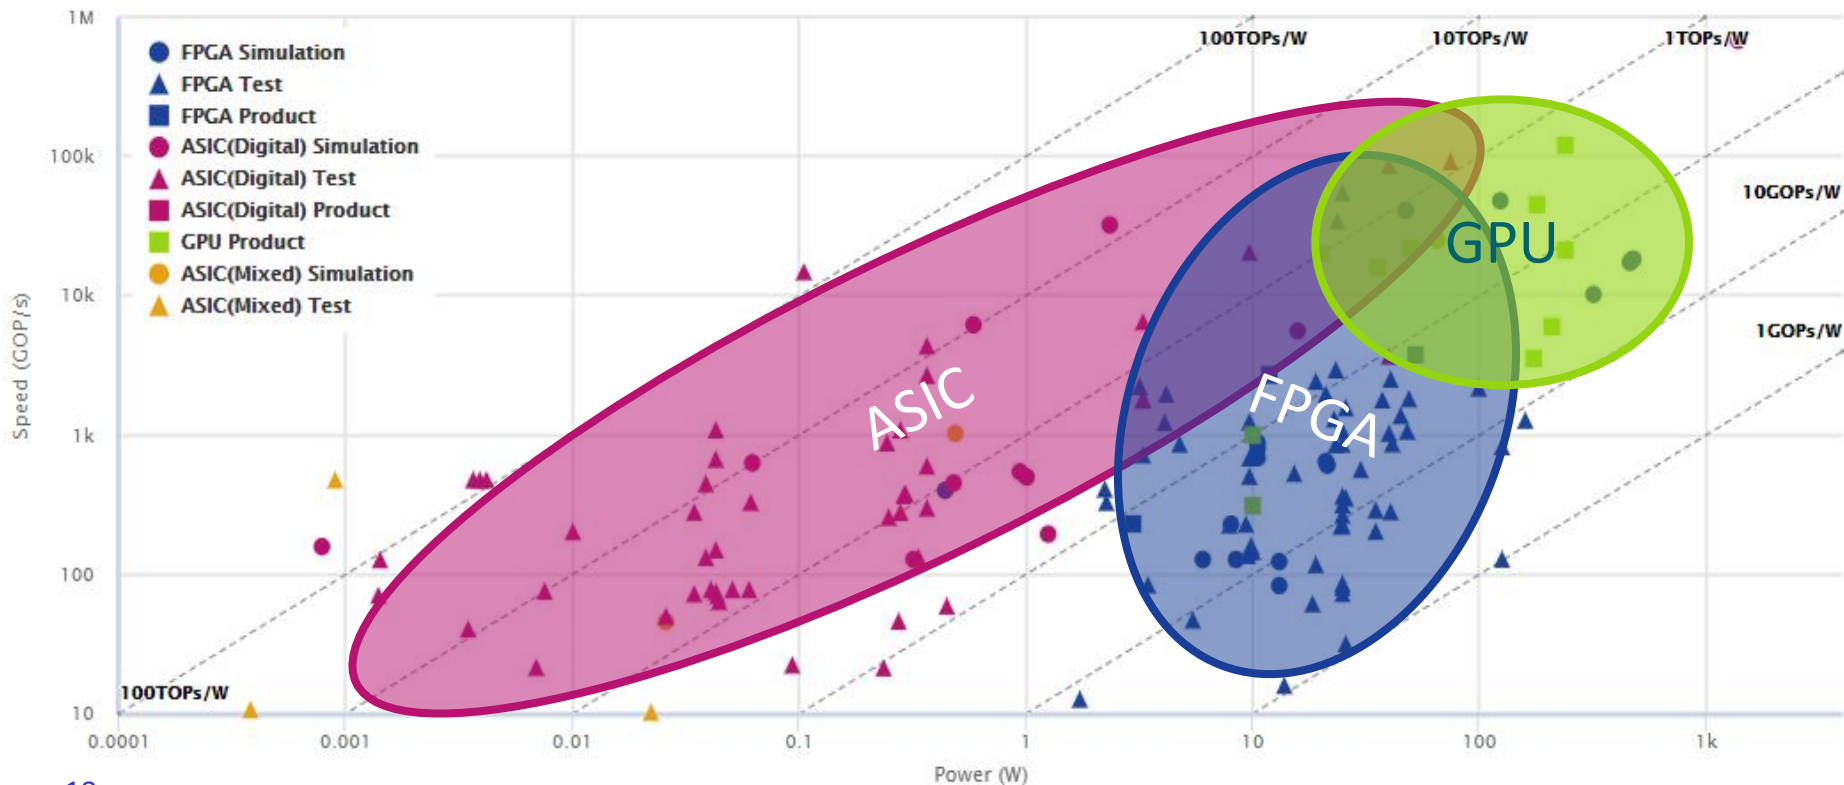


---

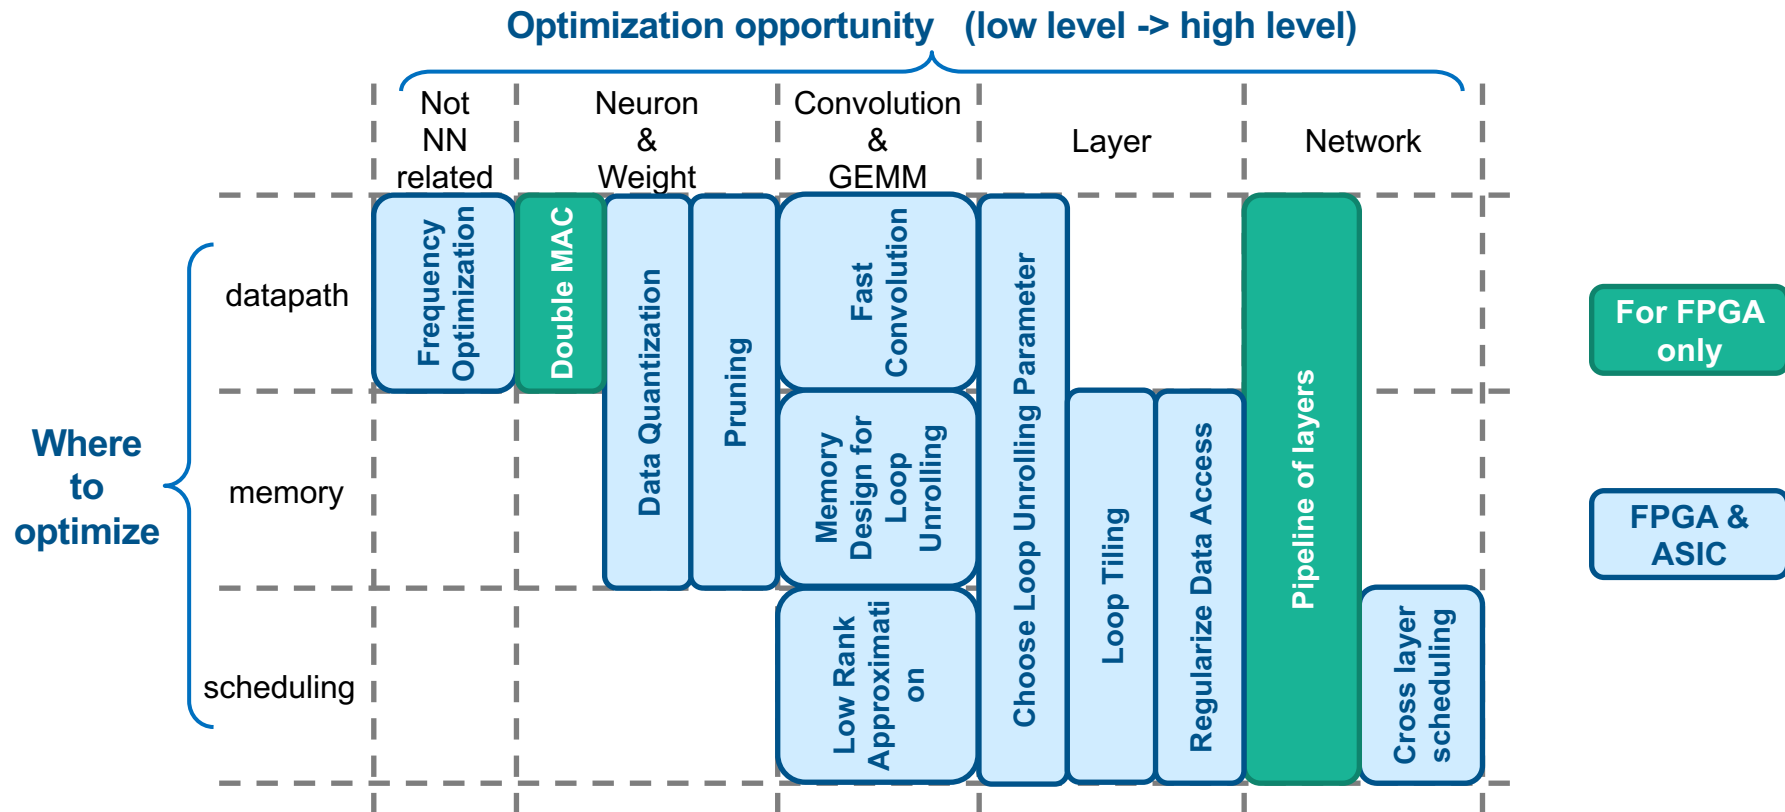
Taking one step back to see the **bigger picture**

# Academic NN Accelerators (Performance vs Power)

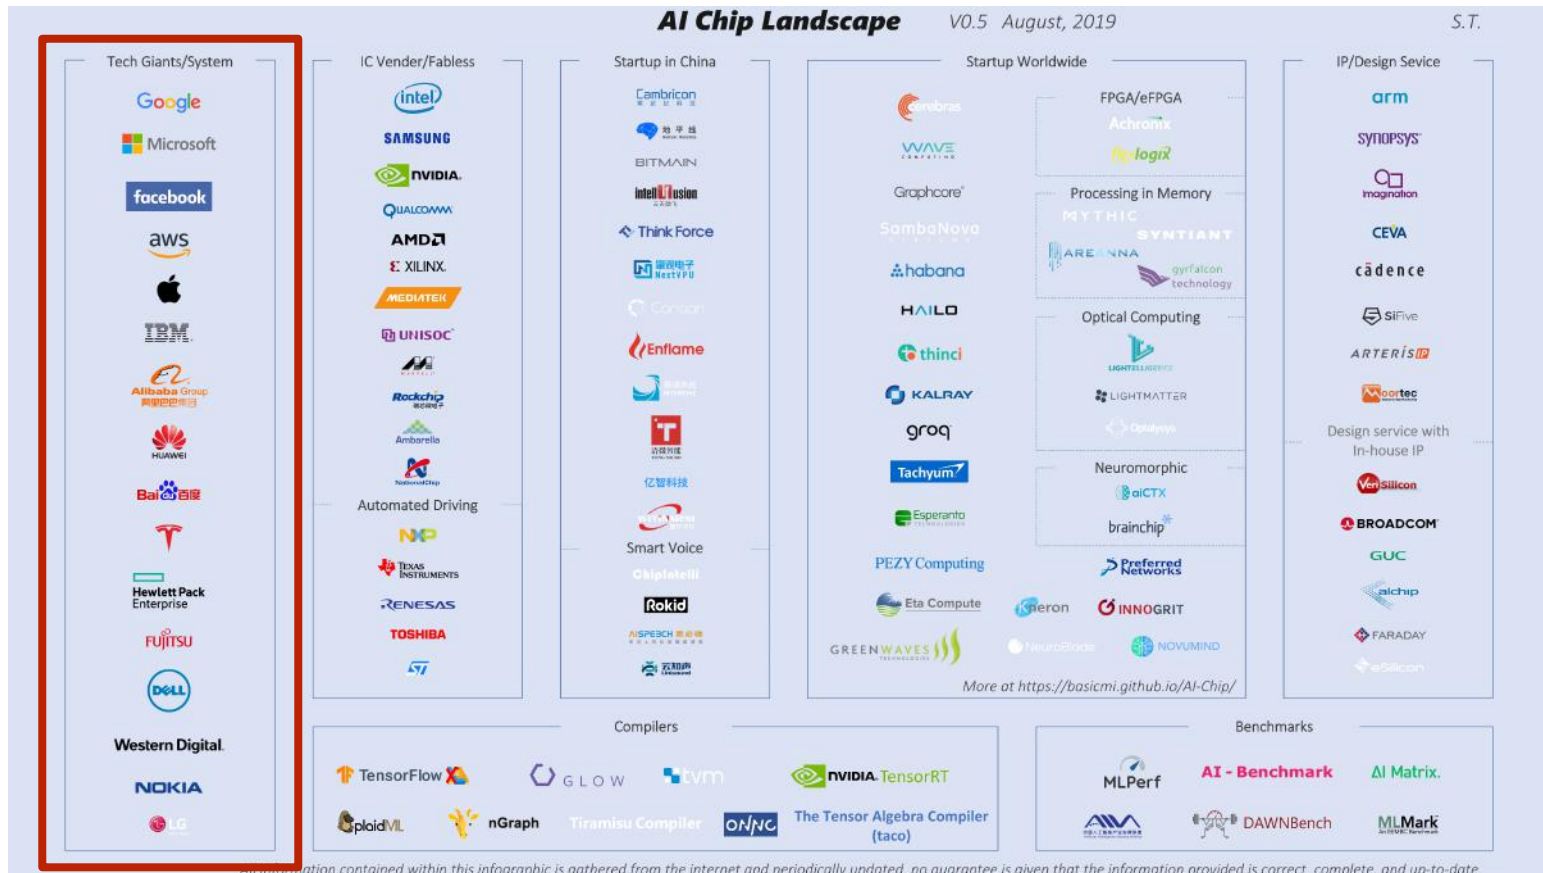
<https://nicsefc.ee.tsinghua.edu.cn/projects/neural-network-accelerator/>



# Survey on FPGA based Inference Accelerators

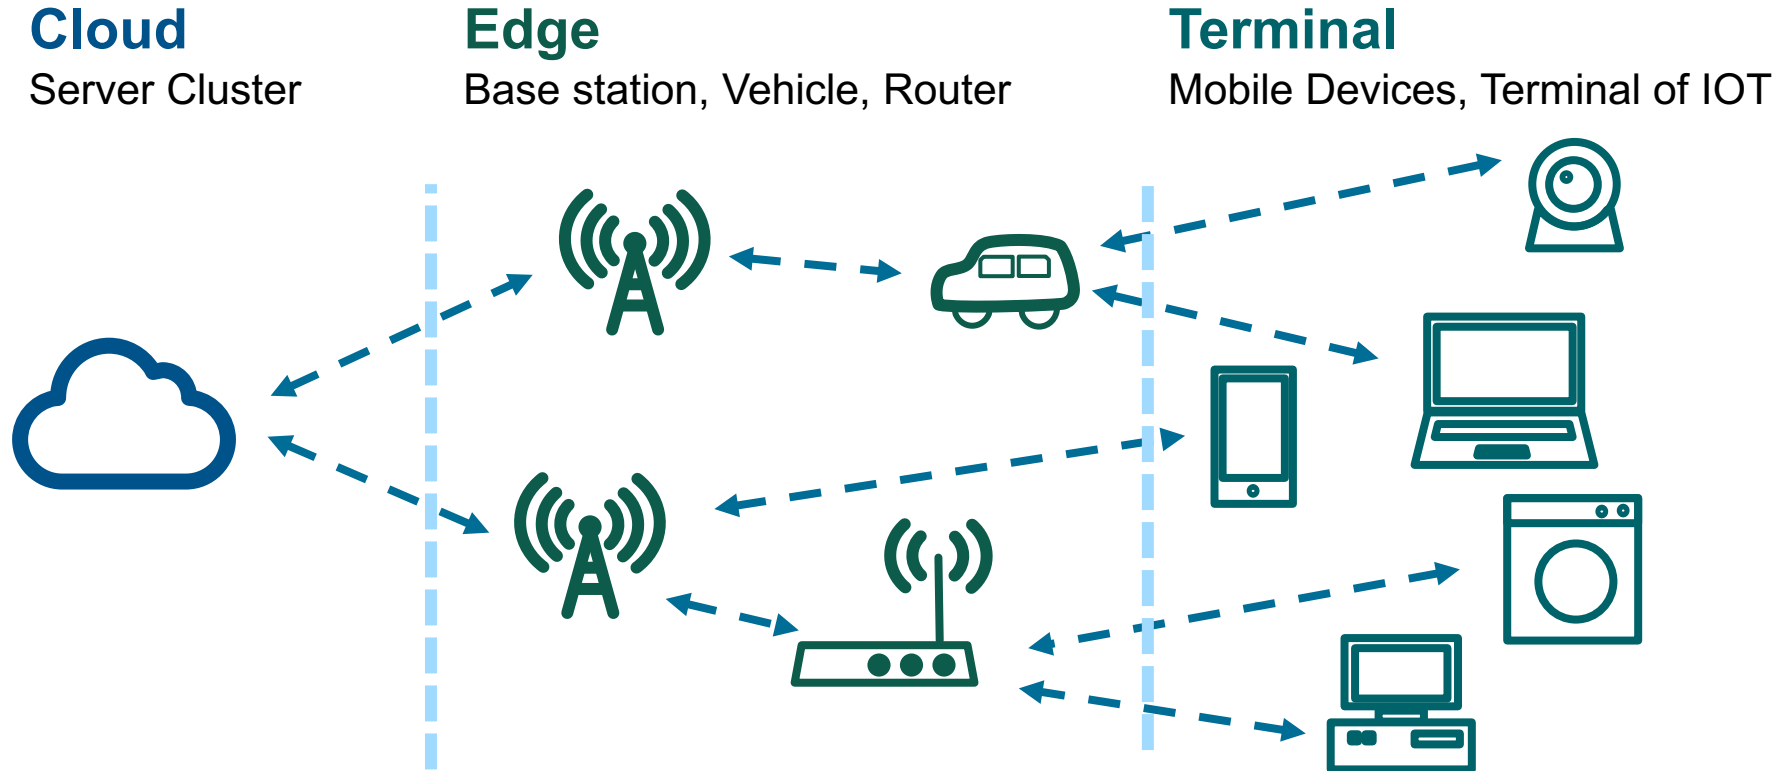


# Industry is Crowding into this Region



All information contained within this infographic is gathered from the internet and periodically updated, no guarantee is given that the information provided is correct, complete, and up-to-date.

# Application Scenarios: Cloud, Edge, Terminal



# Application Scenarios: Cloud, Edge, Terminal



**Cloud**

**Computing + Storage**



**Edge**

**Computing + Communication**



**Terminal**

**Display + Interaction**

**Google  
TPU 3.0**



**Huawei  
Ascend**



**Alibaba  
Hanguang800**



**Xilinx DPU**



**Amba CV2**



**Rockchip**



**Intel  
SpringCrest**



**Xilinx  
ACAP Versal**



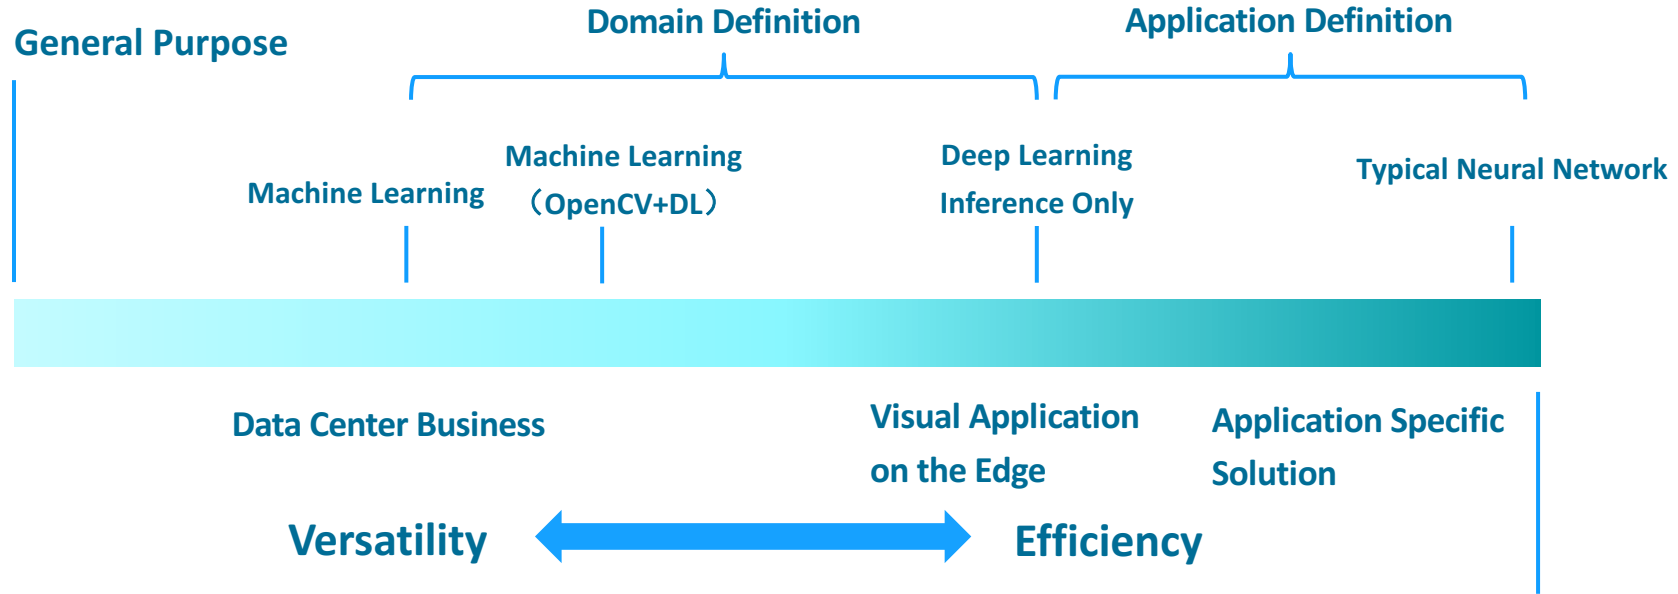
**Nvidia TX2**



**Mobileye EyeQ4**




# Choose Your Business Scope: AI is a Large Area



**Everyone is building their own SOC solution:  
Manufacturing + IP + Tools + Integration Technology**


**Application Specific**

# Growing of Computation Power






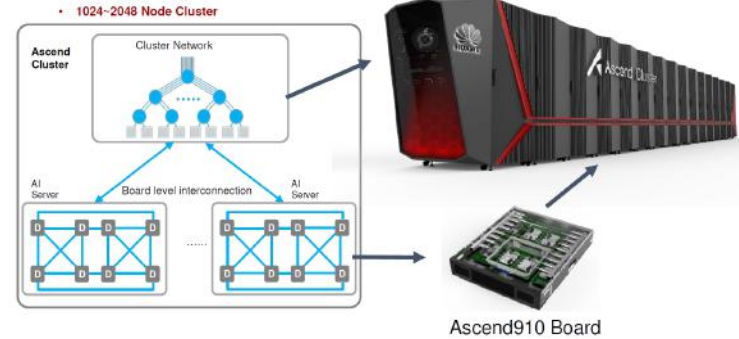
# Even Larger?



**TPU v3 pod: >100PFLOPS**

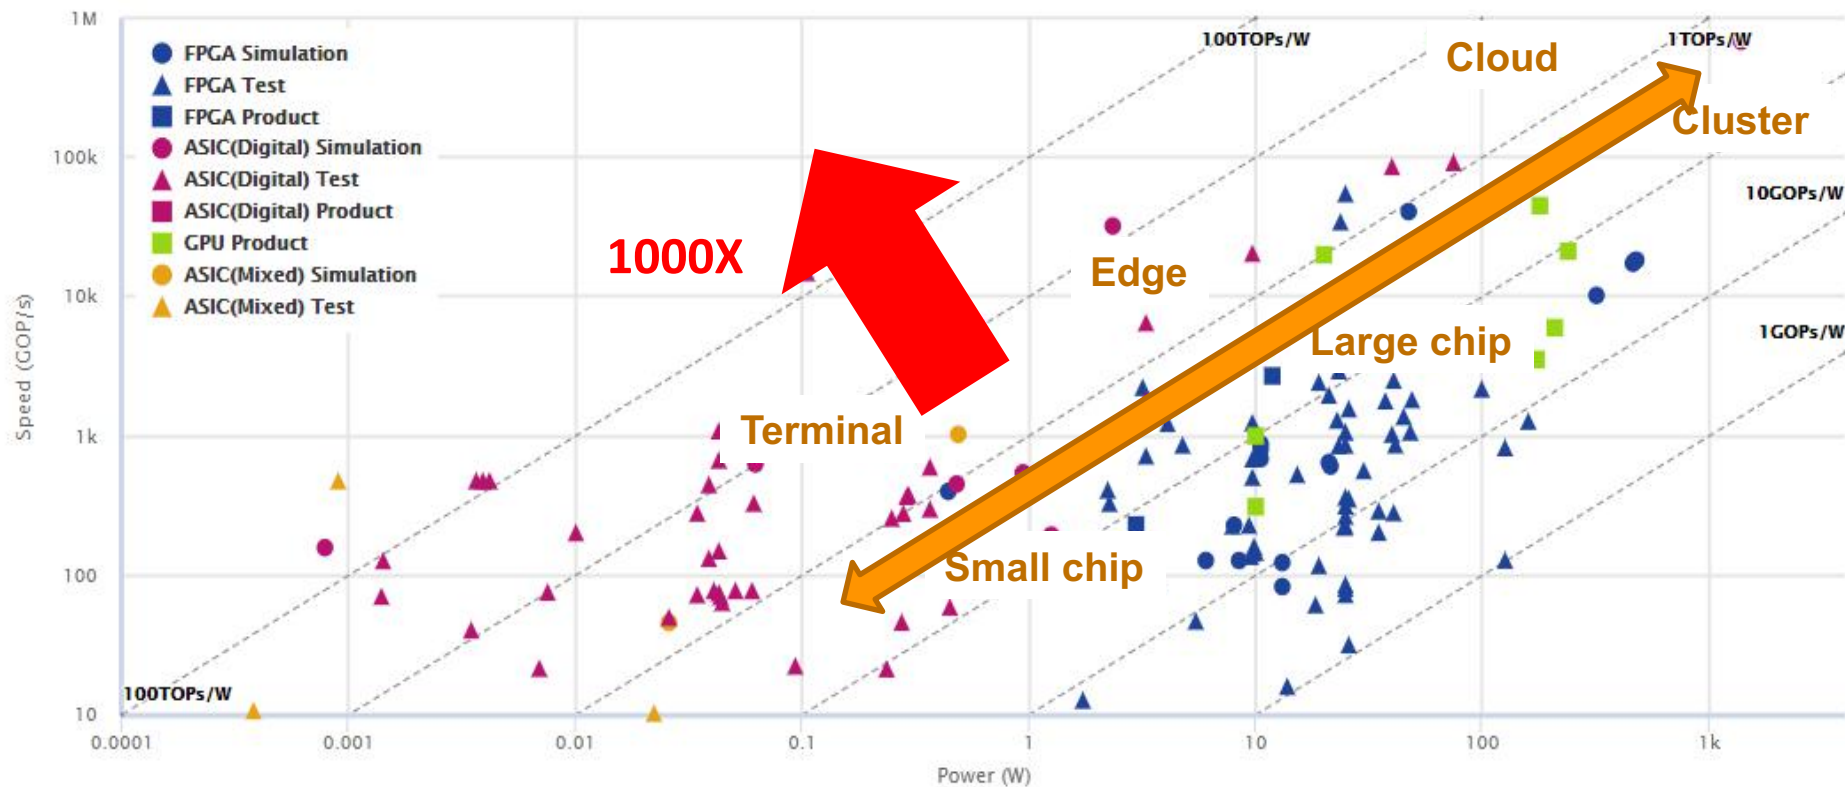


**Cerebras: Wafer-scale chip  
with 1.4T transistors**



**Ascend 910 Cluster: 512PFLOPS**


# We Already Have Many Choices Now



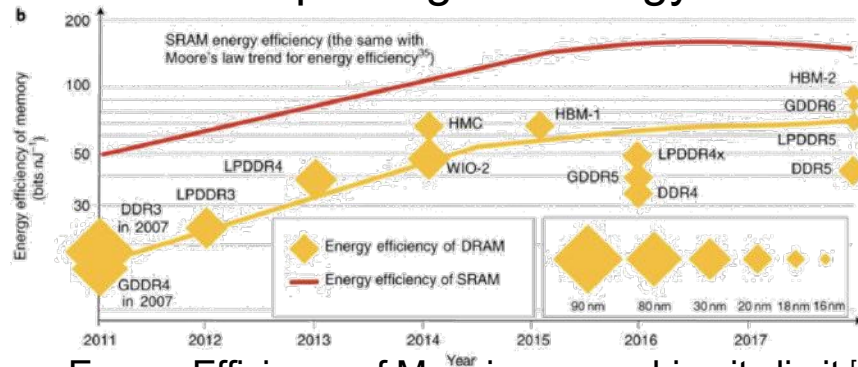
# Bottleneck of Energy Efficiency Improvement

- **Data movements** become the main bottleneck for improving the energy efficiency

Large data movements in AI algorithms -> Aggravated Von Neumann bottleneck



Von Neumann Architecture



Energy Efficiency of Mem is approaching its limit <sup>[1]</sup>

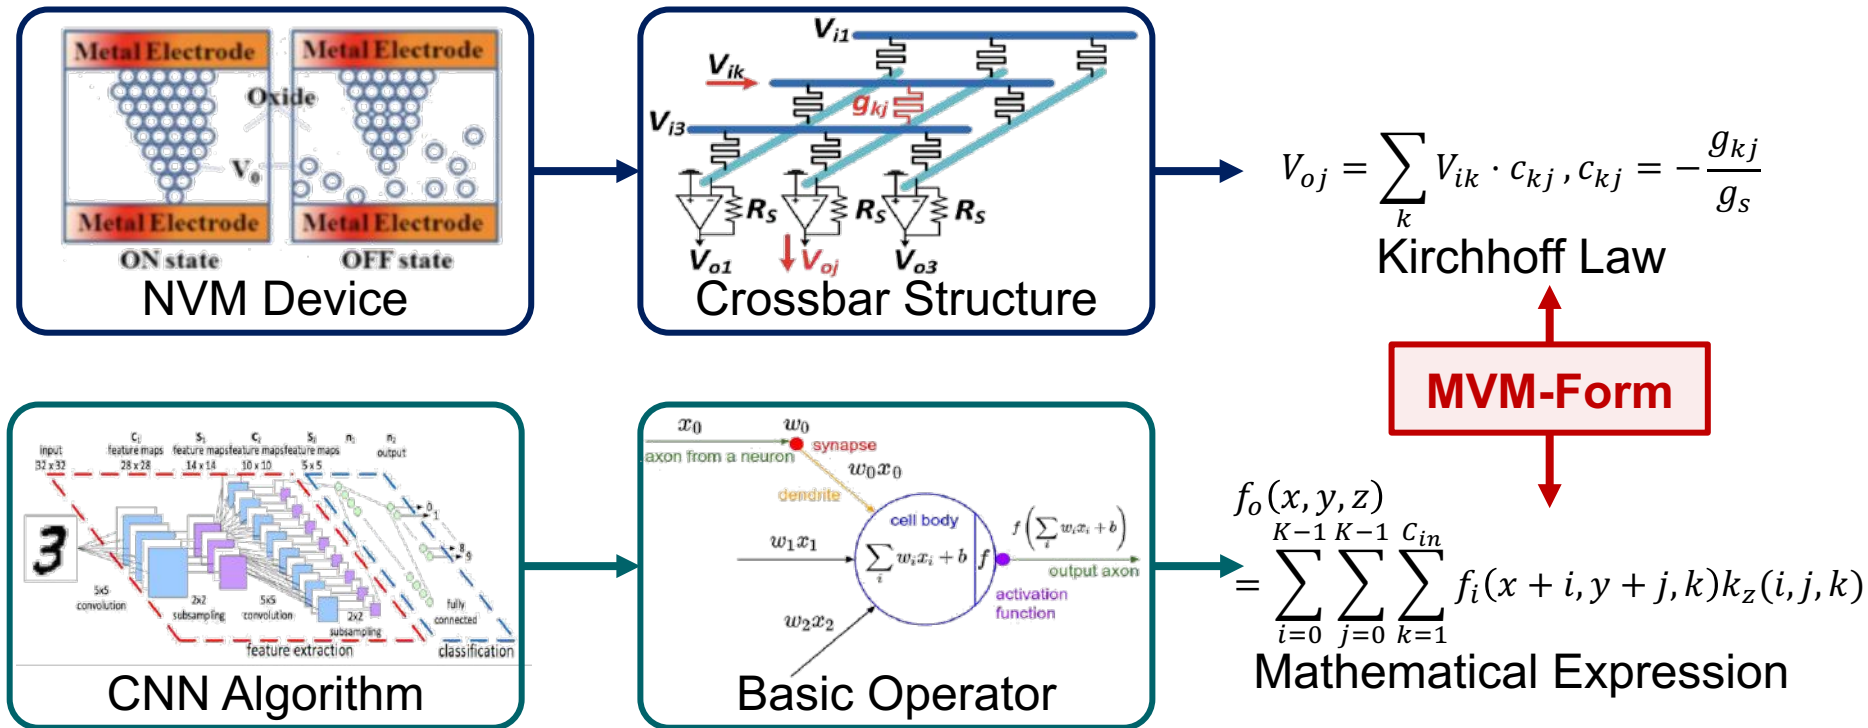
Operation	Energy [pJ]	Relative Cost
32 bit int ADD	0.1	1
32 bit float ADD	0.9	9
32 bit Register File	1	10
32 bit int MULT	3.1	31
32 bit float MULT	3.7	37
32 bit SRAM Cache	5	50
<b>32 bit DRAM Memory</b>	<b>640</b>	<b>6400</b>

High Memory Access Energy Consumption <sup>[2]</sup>

[1] Xu, Xiaowei, et al. "Scaling for edge inference of deep neural networks." Nature Electronics 1.4 (2018): 216.

[2] Han, Song, et al. "Learning both weights and connections for efficient neural network." Advances in NIPS. 2015.



# Non-Volatile Memory (NVM) - based CNN Accelerator



# NVM-based CNN Accelerator

## Academic

- RRAM chip, NTHU, ISSCC19
- 256 rows x 512 columns
- CNN computing: 11.75~14.6ns
- **Energy Efficiency:**
  - **21.9~53.17 TOPS/W**
- CNFET+ RRAM chip, Stanford, ISSCC18
- 1952 CNFETs + 224 RRAM cells
- Pairwise classification of 21 European Languages: 98% accuracy



[1] Xue, Cheng-Xin, et al. "24.1 A 1Mb Multibit ReRAM Computing-In-Memory Macro with 14.6 ns Parallel MAC Computing Time for CNN Based AI Edge Processors." 2019 IEEE ISSCC

[2] Wu, Tony F., et al. "Brain-inspired computing exploiting carbon nanotube FETs and resistive RAM: Hyperdimensional computing case study." 2018 IEEE ISSCC

## Industry

MYTHIC

- Based on embedded NOR FLASH
- Highest energy efficiency: 4TOPS/W
- <https://www.mythic-ai.com/technology/>



- Based on NOR Flash

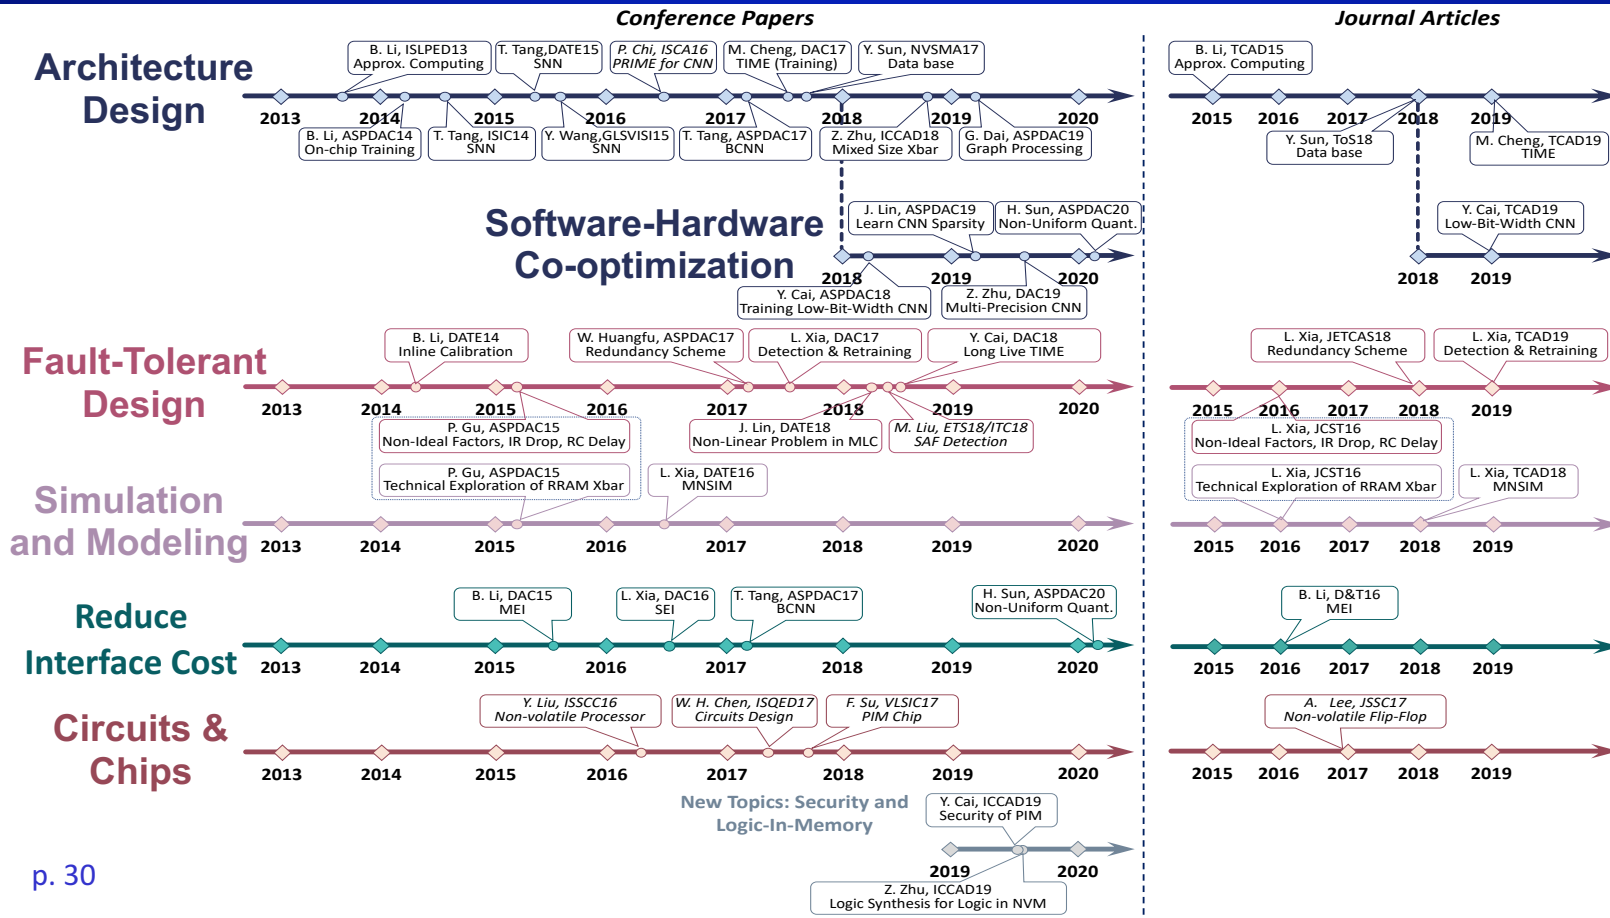
SYNTIANT

- Based on NOR Flash
- Built for voice applications

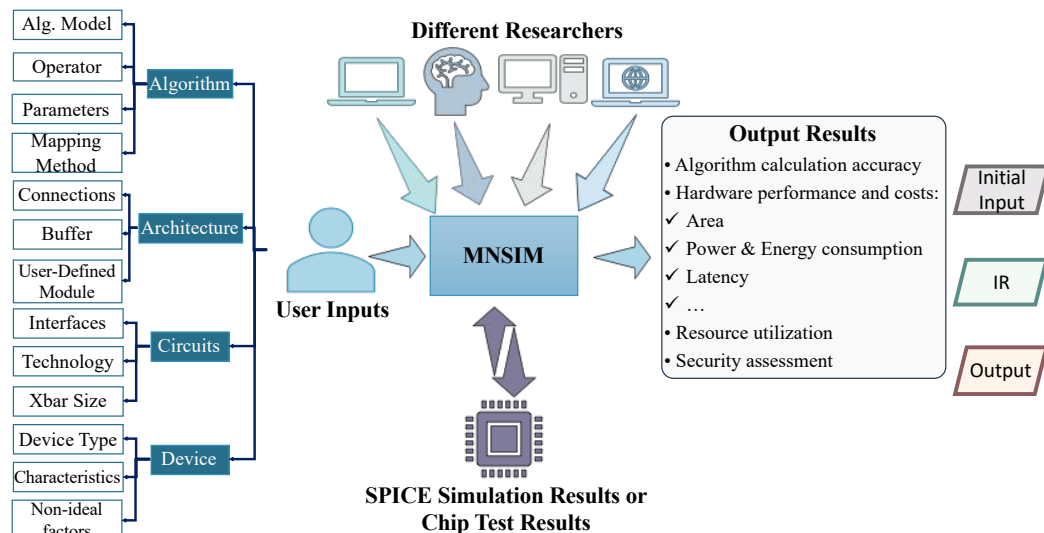


- Based on PCM
- 8-bit precision device (still in research stage)
- IEDM 2018

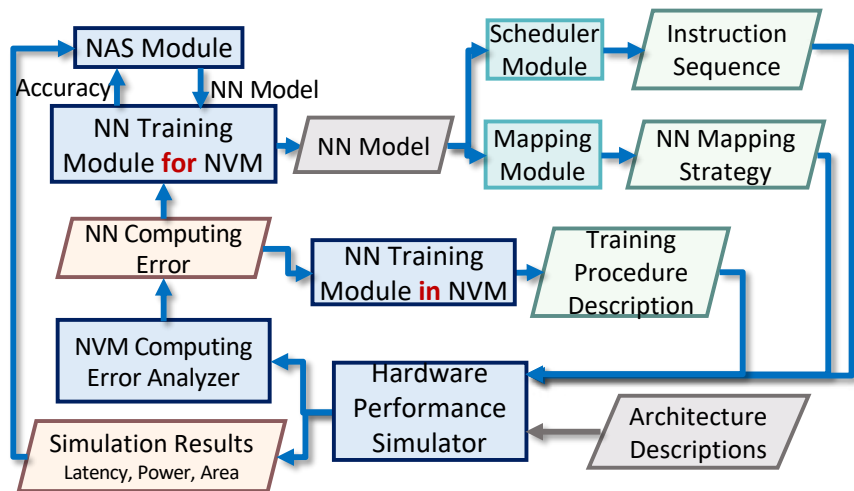
# Research on RRAM-based PIM



# MNSIM: High-Speed Behavior-Level Model for Design Optimization



The overview of MNSIM



The entire modeling flow

**The code is available in:**  
[https://github.com/Zhu-Zhenhua/MNSIM\\_Python](https://github.com/Zhu-Zhenhua/MNSIM_Python)



# Brief Summary

---

## What we have or we can expect

- AI platforms are covering everywhere.
- We already have many methods to make the chips powerful.

## What's next: **how they will influence the hardware and software design for NN on chip**


- Are the chips ready for real applications?
- Can we use the large 'TOPs' to the extreme?
- Are the systems robust enough?
- Are the systems secure enough?





# Case Study 1: Using CNN accelerator in moving robots

# CNN Greatly Benefits Basic Functions in Robotic Applications



## Feature-Point Extraction

SuperPoint <sup>[1]</sup>:

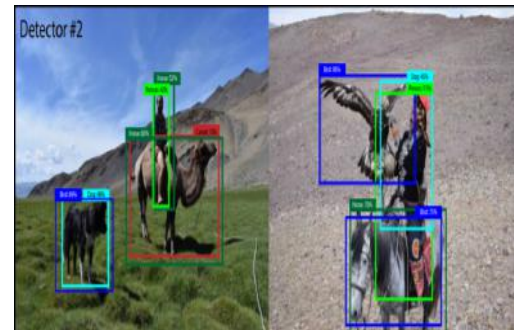
The matching accuracy is **20%** higher than handcrafted method.



## Place Recognition

GeM <sup>[2]</sup>:

The recall of image query is **10%** higher than handcrafted method.



## Object Detection

YOLOv3 <sup>[3]</sup>:

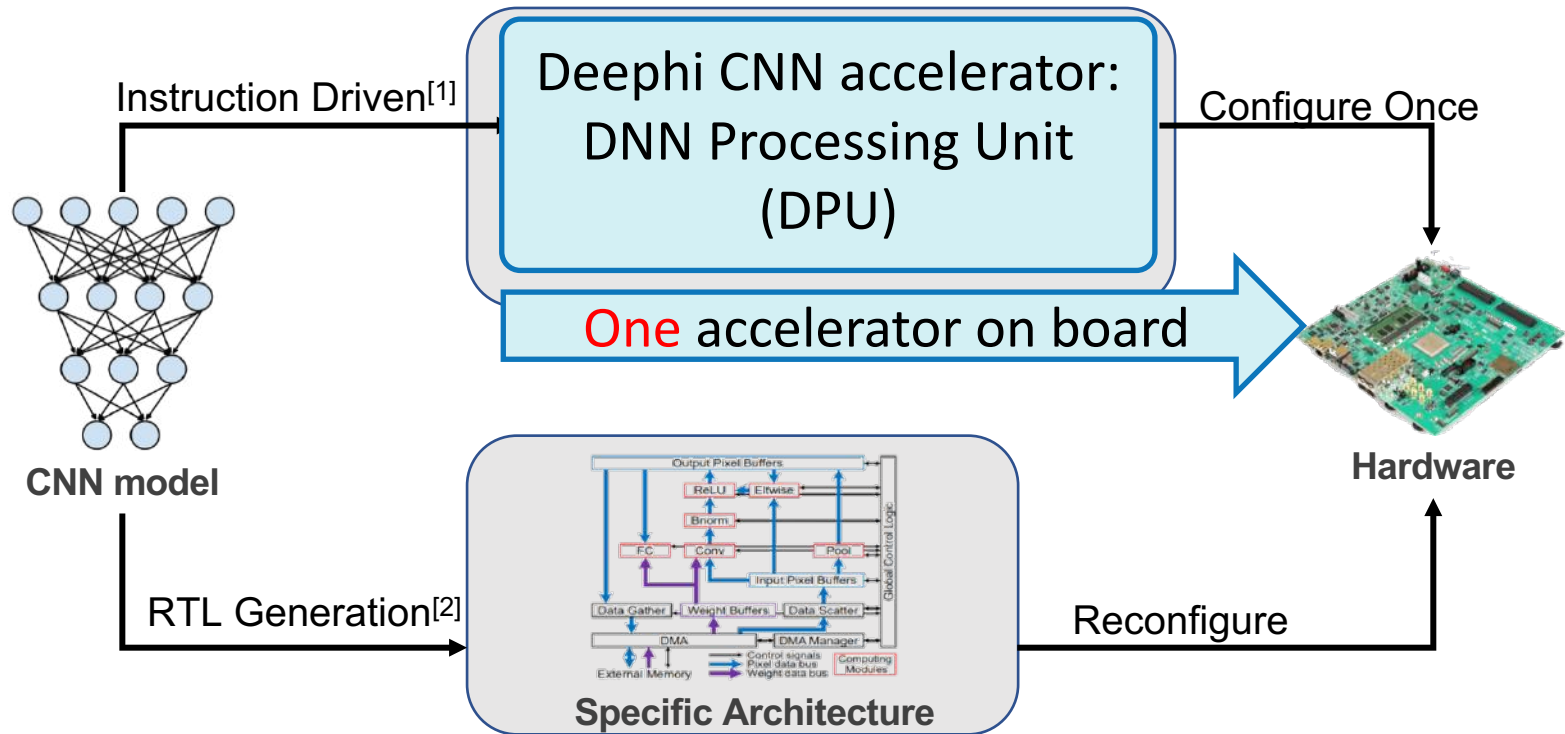
The mean average precision (mAP) on COCO is above **58%**.

[1] Daniel DeTone, Tomasz Malisiewicz, and Andrew Rabinovich. 2018. Superpoint: Self-supervised interest point detection and description. In CVPR Workshops. 224–236.

[2] Filip Radenović, Giorgos Tolias, and Ondřej Chum. 2018. Fine-tuning CNN image retrieval with no human annotation. IEEE transactions on pattern analysis and machine intelligence 41, 7 (2018), 1655–1668.

[3] Redmon, J., & Farhadi, A. (2018). YOLOv3: an incremental improvement. arXiv. 1804.02767.

# Accelerators Enables CNN on Moving Robots




[1] [FPT 17] Jincheng Yu, Yiming Hu, Xuefei Ning, Jiantao Qiu, Kaiyuan Guo, Yu Wang, and Huazhong Yang. 2017. Instruction driven cross-layer CNN accelerator with winograd transformation on FPGA. In International Conference on Field Programmable Technology. 227–230


[2] Yufei Ma, Yu Cao, Sarma Vrudhula, and Jae-sun Seo. 2017. An automatic RTL compiler for high-throughput FPGA implementation of diverse deep convolutional neural networks. In Field Programmable Logic and Applications (FPL), 2017 27th International Conference on. IEEE, 1–8

# Tasks with Hard Deadline

**Visual Odometry (VO):**  
Basic for localization & obstacle detection




Obstacle detected on time :  
Circumvention



Obstacle detected not on time :  
Collision

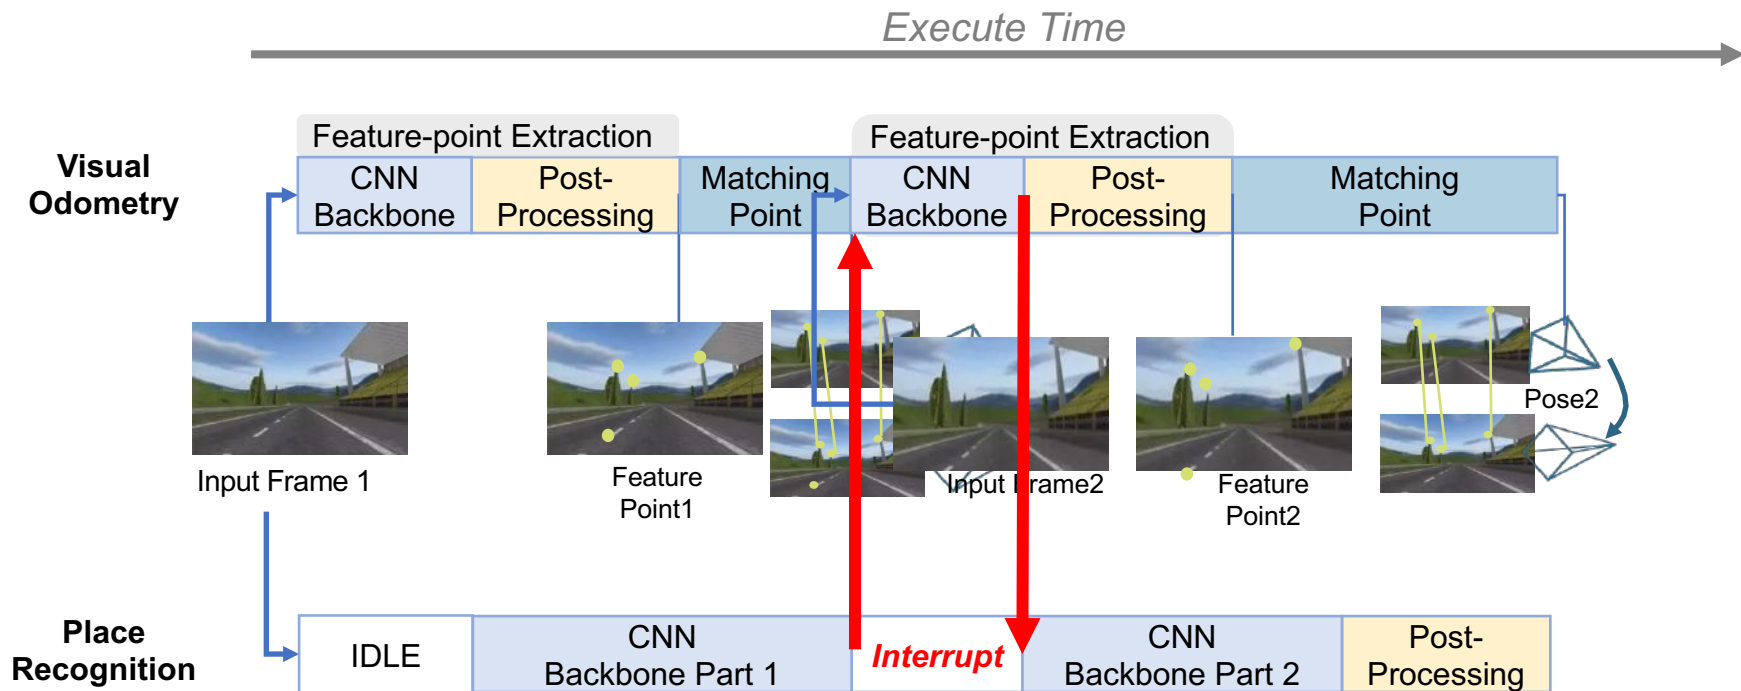
**Place Recognition (PR):**  
For map merging & optimization



Map merging can be delayed

**Some tasks need to be finished before **deadline****

# Accelerator Interrupt for Hardware Conflicts



If we only have **ONE** DPU, how to share between two threads?

# Interrupt Respond Latency & Extra Cost

## • 4 Stages for Interrupt

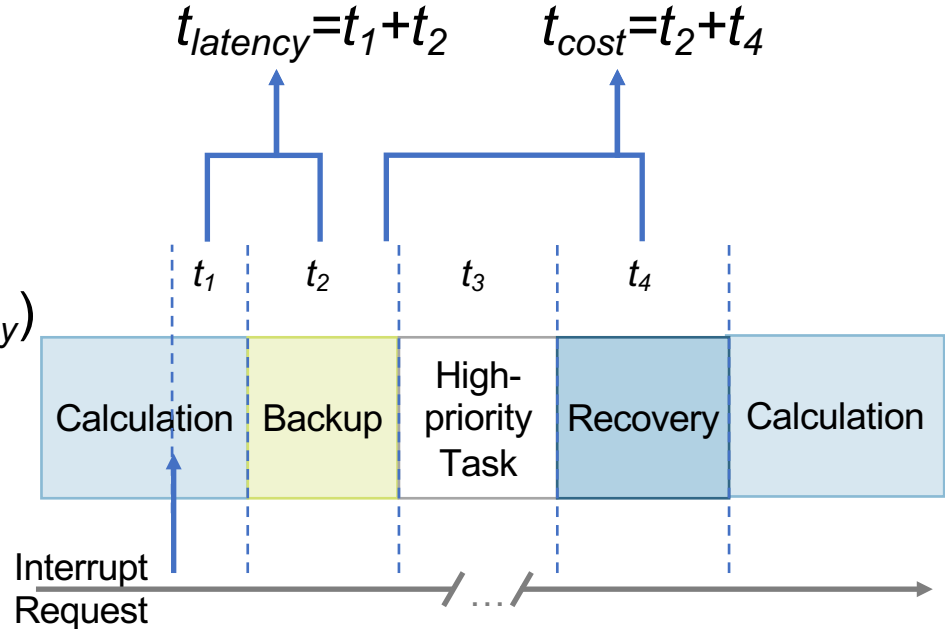
- To finish current instruction ( $t_1$ )
- To backup low-priority task ( $t_2$ )
- To execute high-priority task ( $t_3$ )
- To recovery low-priority task ( $t_4$ )

## • Interrupt Respond Latency ( $t_{latency}$ )

- $t_{latency} = t_1 + t_2$

## • Extra Cost ( $t_{cost}$ )

- $t_{cost} = t_2 + t_4$

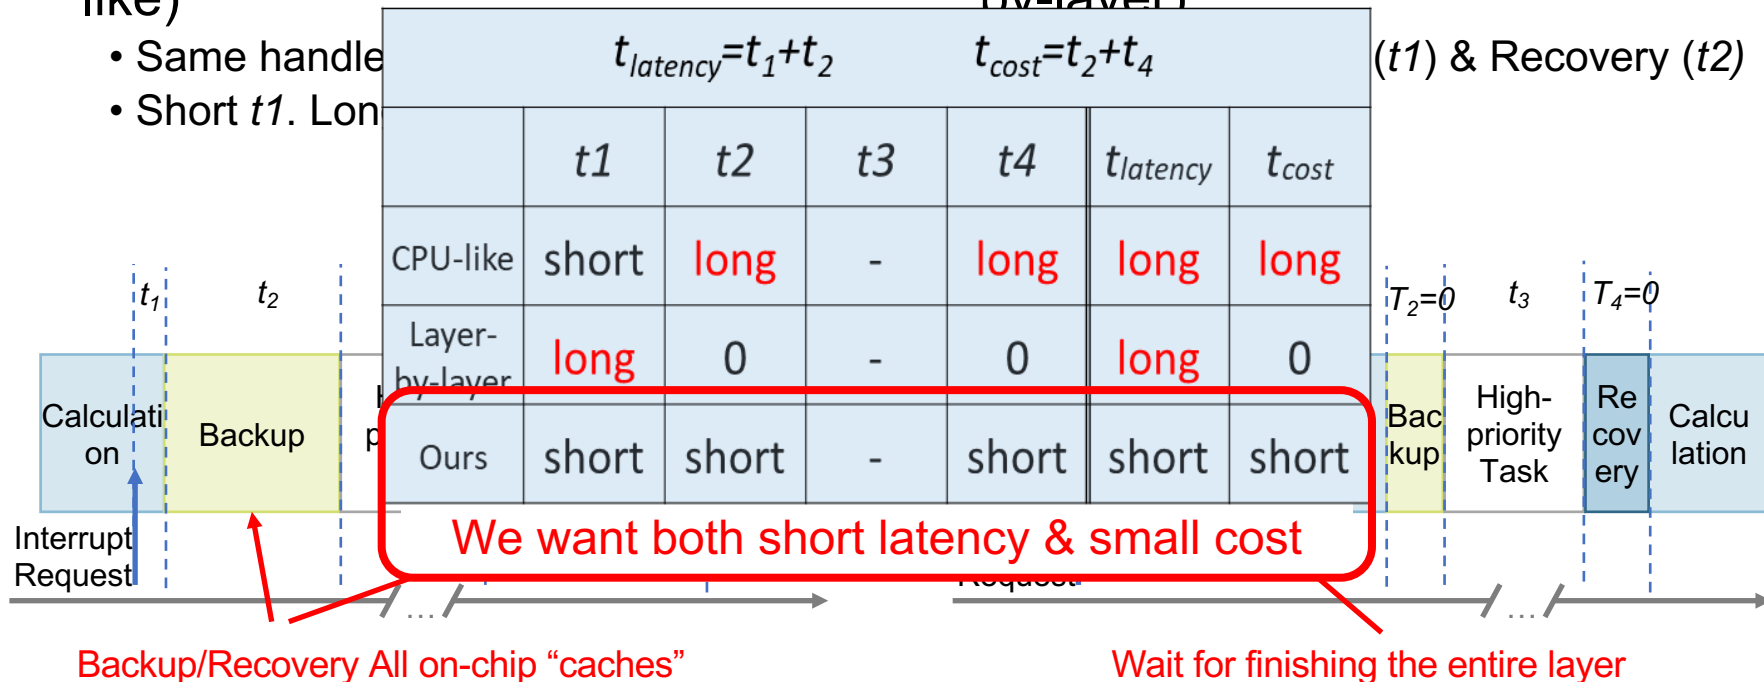


# How to Interrupt?

- Interrupt Intermediately (CPU-like)

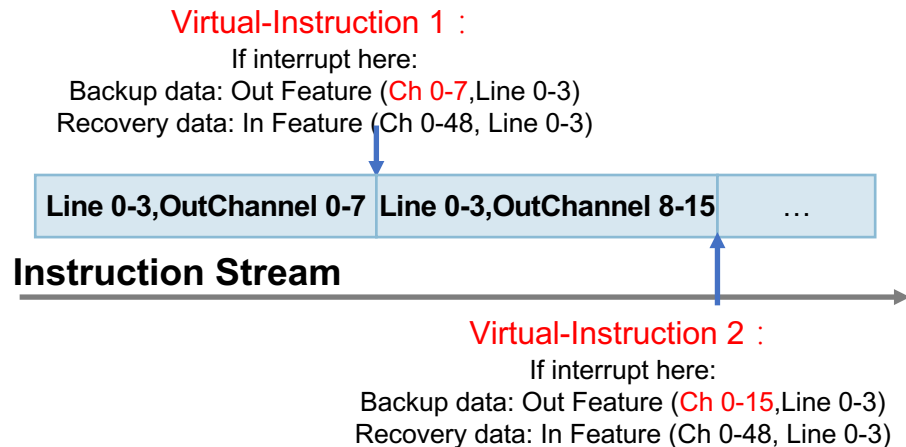

- Same handle
- Short  $t_1$ . Long

- Interrupt After Each layer (Layer-by-layer)



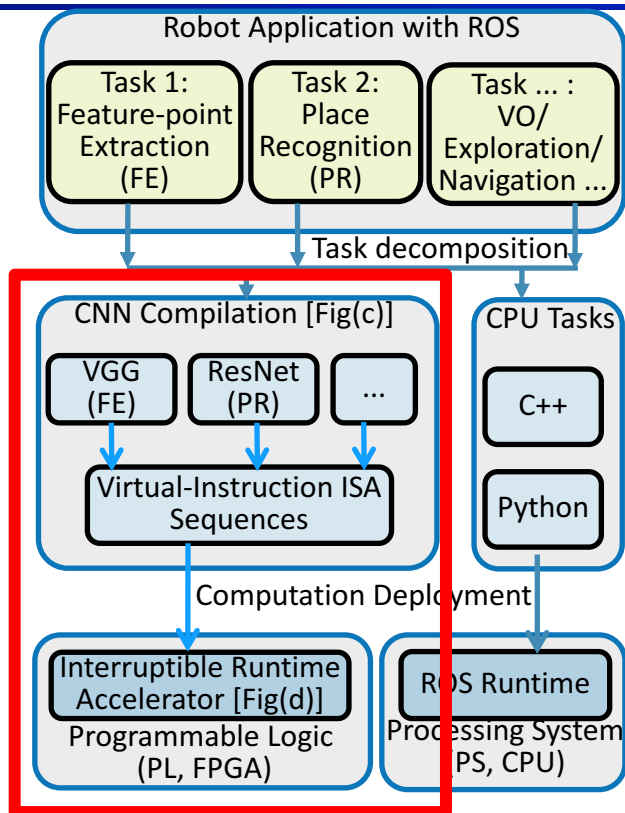
# Virtual-Instruction-Based Interrupt

- Interruptible at some positions with small extra cost
- Different handler for different interrupt position
  - Some Virtual-instructions are inserted into the instruction sequence.

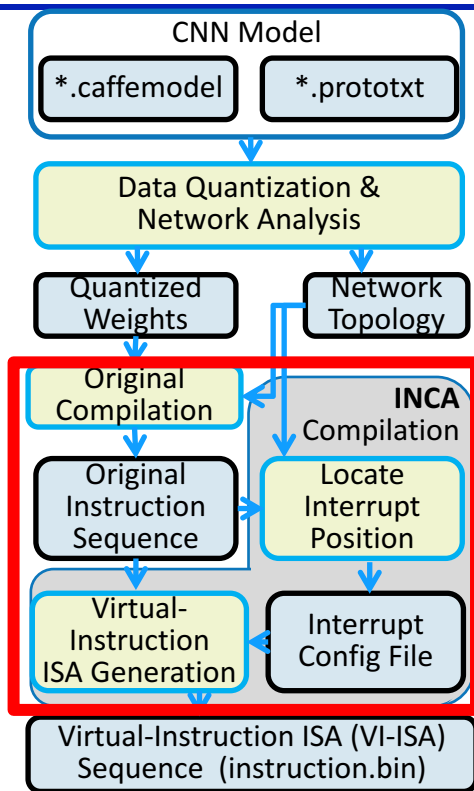




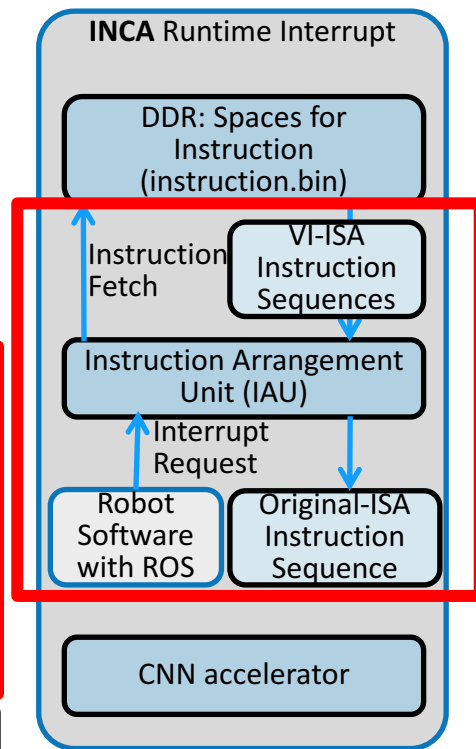
# INCA: Mapping Robot Applications to FPGA



(a) INCA framework from robot application to hardware platform.



(b) At compilation, INCA adds virtual instructions to produce VI-ISA instruction sequences.

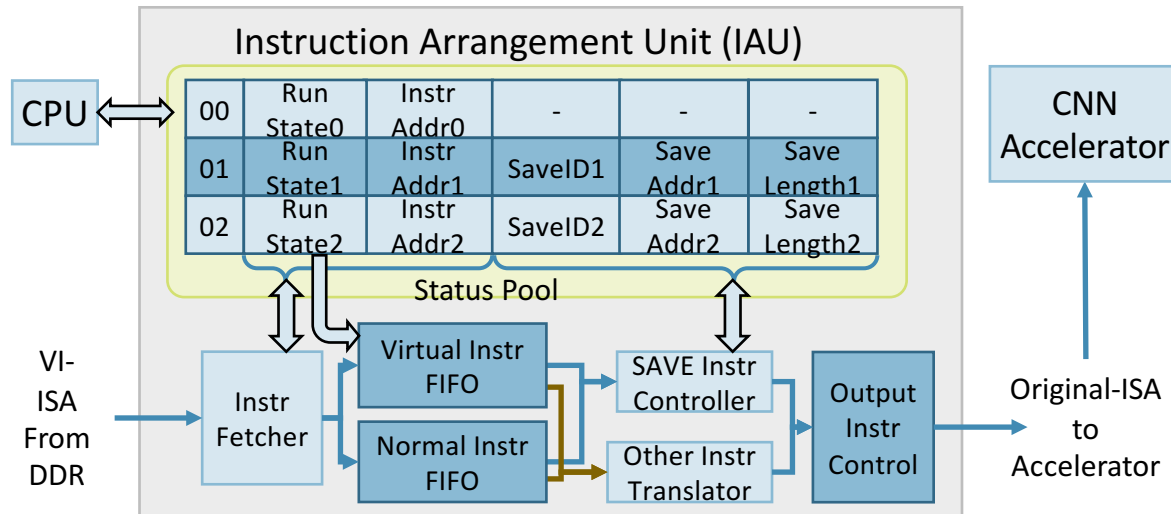


(c) At runtime, INCA translates the VI-ISA to original ISA executed on the accelerator.

# Instruction Arrangement Unit (IAU)

- IAU

- Monitor interrupt status
- Fetch instructions for different tasks
- Translate VI-ISA to original ISA



# Example of VI-ISA

	Type	Address1	Address2	Address3	Workload	Virtual	SaveID
1	LOAD_D	0x0000	0x4000	-	0x20	2'b00	1
2	LOAD_W	0x1000	0x5000	-	0x1	2'b00	1
3	CALC_F	0x4000	0x6000	0x5000	0x20	2'b00	1
4	SAVE	0x2000	0x6000	-	0x20	2'b01	1
5	LOAD_D	0x0000	0x4000	-	0x20	2'b10	1
6	LOAD_W	0x1001	0x5000	-	0x1	2'b00	1
7	CALC_F	0x4000	0x6020	0x5000	0x20	2'b00	1
8	SAVE	0x2000	0x6000	-	0x40	2'b00	1

(a) Input instruction sequence



	Type	Address1	Address2	Address3	Workload
1	LOAD_D	0x0000	0x4000	-	0x20
2	LOAD_W	0x1000	0x5000	-	0x1
3	CALC_F	0x4000	0x6000	0x5000	0x20
4	LOAD_W	0x1001	0x5000	-	0x1
5	CALC_F	0x4000	0x6020	0x5000	0x20
6	SAVE	0x2000	0x6000	-	0x40

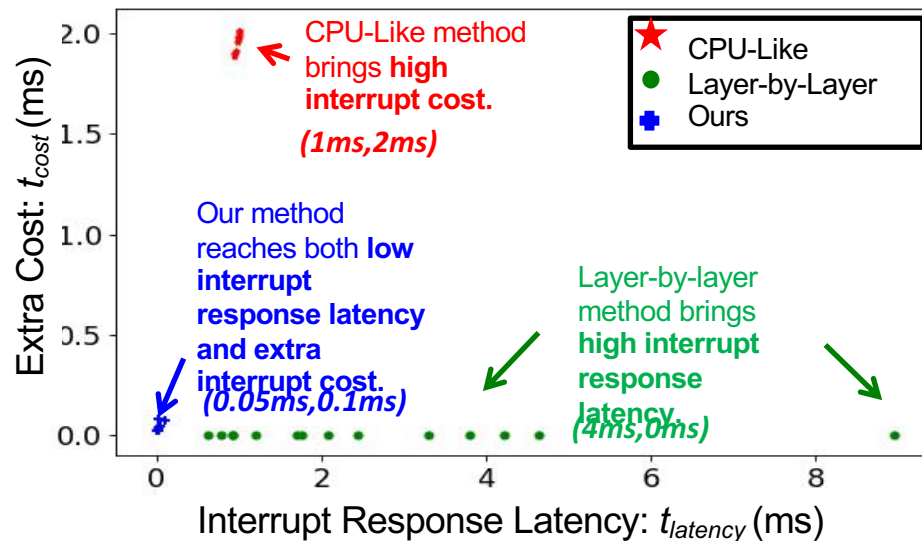



(b) Executed sequence when no interrupt.  
Virtual Instructions are deleted.

	Type	Address1	Address2	Address3	Workload
1	LOAD_W	0x1000	0x5000	-	0x1
2	LOAD_D	0x0000	0x4000	-	0x20
3	CALC_F	0x4000	0x6000	0x5000	0x20
4	SAVE	0x2000	0x6000	-	0x20
HIGH-PRIORITY CNN					
5	LOAD_D	0x0000	0x4000	-	0x20
6	LOAD_W	0x1001	0x5000	-	0x1
7	CALC_F	0x4000	0x6020	0x5000	0x20
8	SAVE	0x2020	0x6020	-	0x20

(c) Executed sequence when interrupt occurs.  
Virtual Instructions (Blue) are executed.  
Normal SAVE (Red) is modified.


# Results for Interrupt ( $T_{cost}$ and $T_{latency}$ )

- Low-priority interruptible CNN backbone: Example on ResNet-101



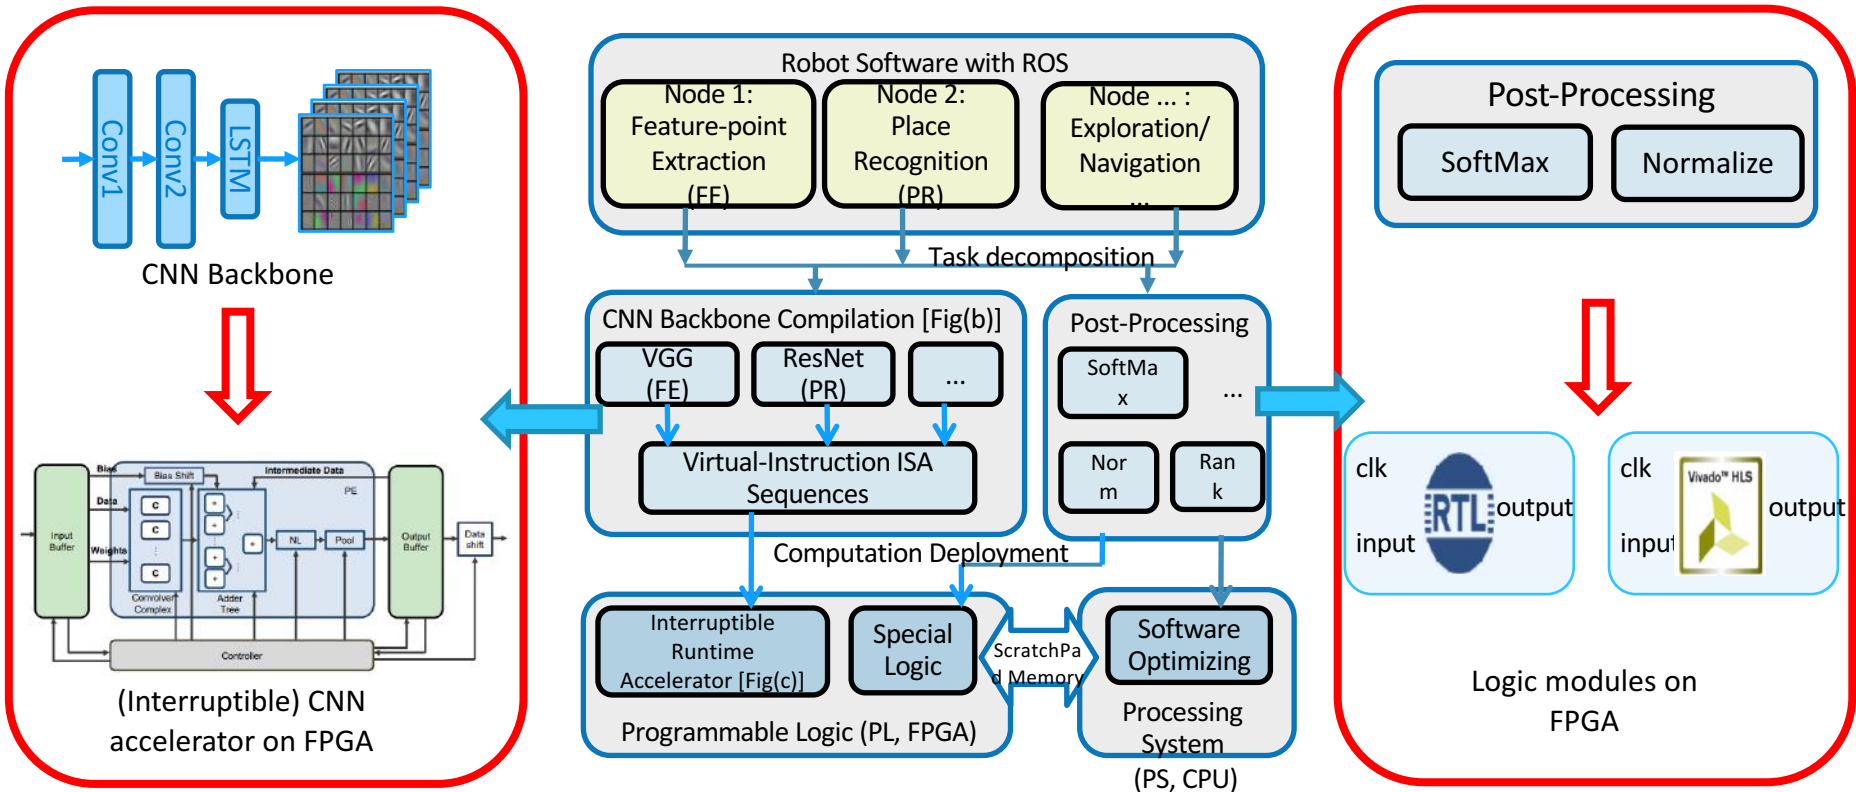
# Latency Comparison ( $T_{\text{latency}}$ )

- Test on different CNN networks: ~100us latency



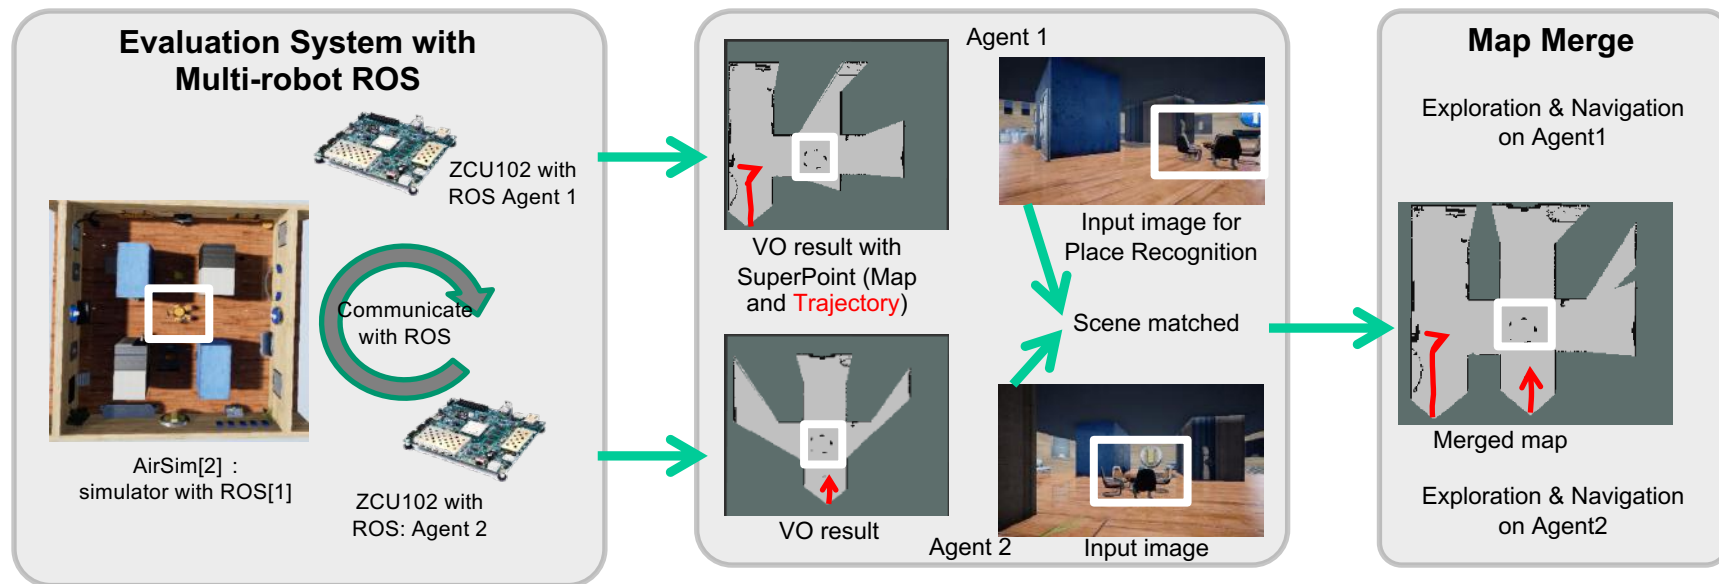
Layer	A	B	C	D	E	F	G	H	I
Network	ResNet101			VGGNet			MobileNet		
W x H	41 x 31	160 x 120	640 x 480	7 x 7	56 x 56	224 x 224	14 x 14	56 x 56	112 x 112
Chin x Chout	1024 x 256	128 x 128	3 x 64	512 x 512	256 x 256	64 x 128	512 x 512	256 x 256	32 x 32
kernel	1	3	7	3	3	3	1	1	3

# Map Robot Applications onto FPGA



# Multi-robot Exploration/SLAM with Hardware-In-Loop Simulation

- Multi-robot Exploration with ROS and Interruptible DPU on Xilinx Board
  - Feature-point extraction 20fps, place recognition (loop-closure detection) 2 fps.



[1] Morgan Quigley, Ken Conley, Brian Gerkey, Josh Faust, Tully Foote, Jeremy Leibs, Rob Wheeler, and Andrew Y Ng. 2009. ROS: an open-source Robot Operating System. In ICRA workshop, Vol. 3. Kobe, Japan, 5.

[2] Shital Shah, Debadeepta Dey, Chris Lovett, and Ashish Kapoor. 2018. Airsim: High-fidelity visual and physical simulation for autonomous vehicles. In Field and service robotics. Springer, 621–635.



# Case Study 2:


## Can we use the large 'TOPs' to the extreme?

### -- **Sharing on the Cloud**



# DNN Inference Tasks in the Cloud

- Deep neural network (DNN) inference tasks take up the majority of the total deep learning workloads in data centers.



DL inference tasks is **doubling each year** in Facebook data center <sup>[1]</sup>




Inference tasks make up nearly **90%** of the total deep learning tasks in Amazon data center <sup>[2]</sup>

[1] Jongsoo Park et al. 2018. Deep learning inference in facebook data centers: Characterization, performance optimizations and hardware implications. arXiv.


[2] Andy Jassy. 2018. Amazon AWS ReInvent Keynote. <https://www.youtube.com/watch?v=ZOIkOnW640A>.

# FPGAs for DNN Inference

- FPGAs are promising acceleration platforms for DNN inference in the cloud[1].

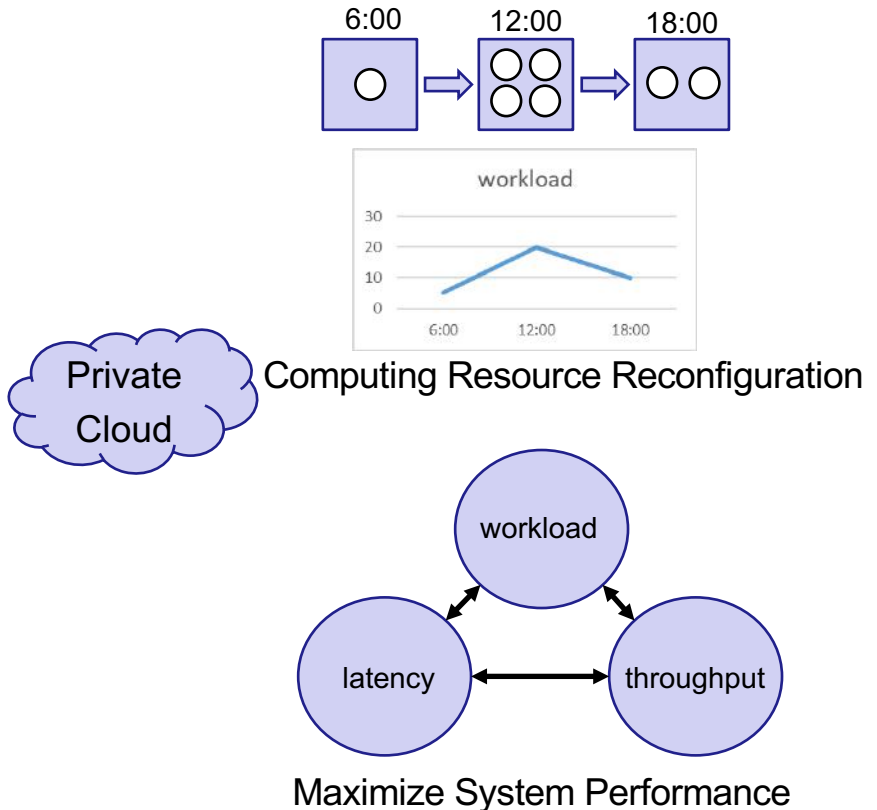
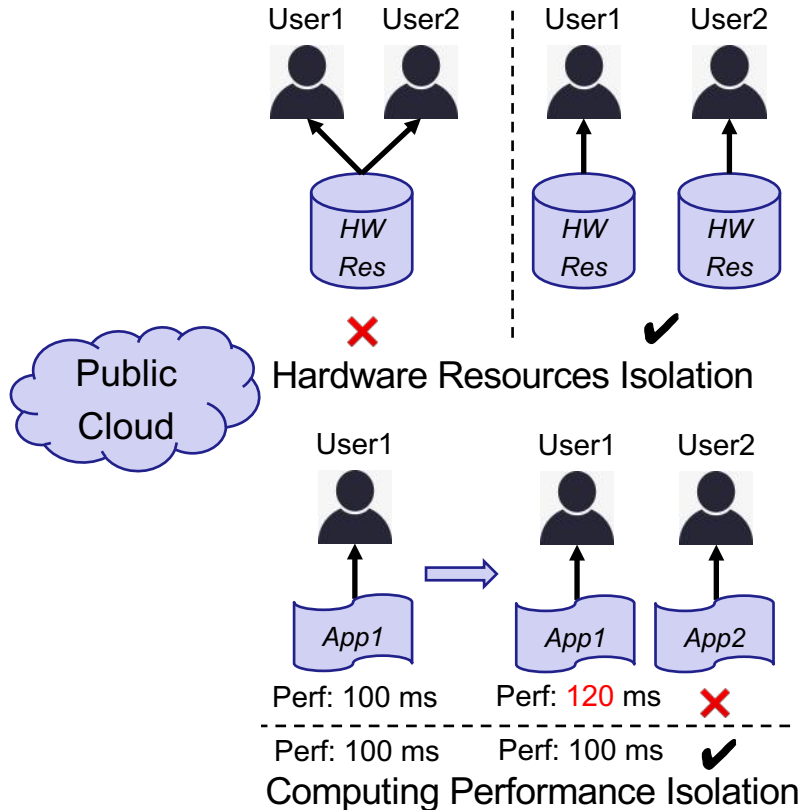


Throughput and Latency

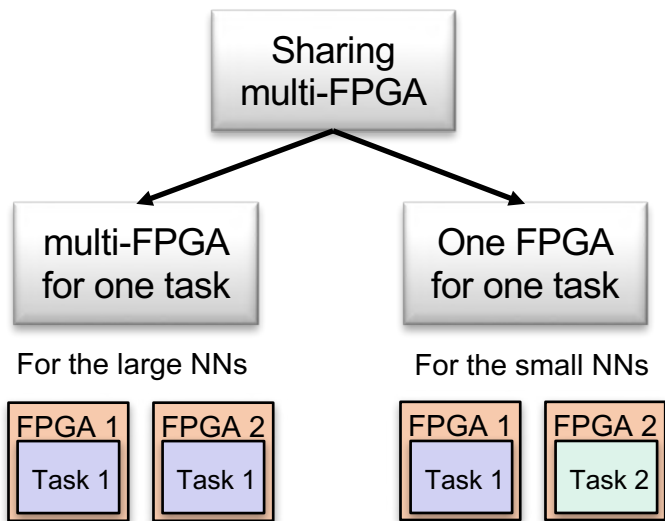


Energy Efficiency


# Key Ideas of Virtualization



# FPGA Sharing in the Cloud

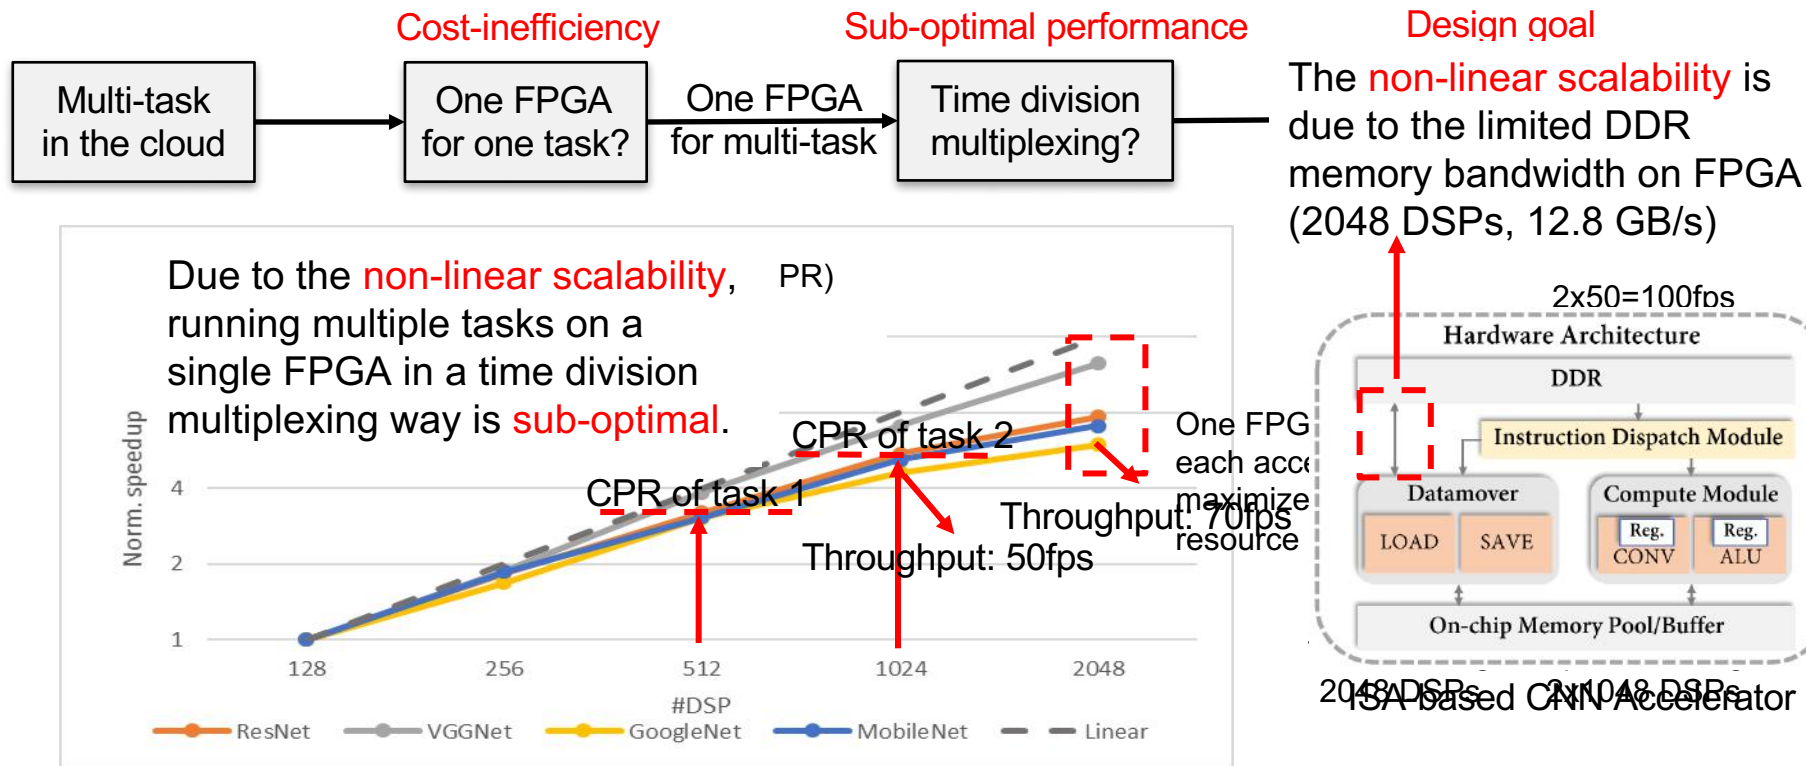


- Good physical and performance isolation
- High reconfiguration overhead
- Cost-inefficient



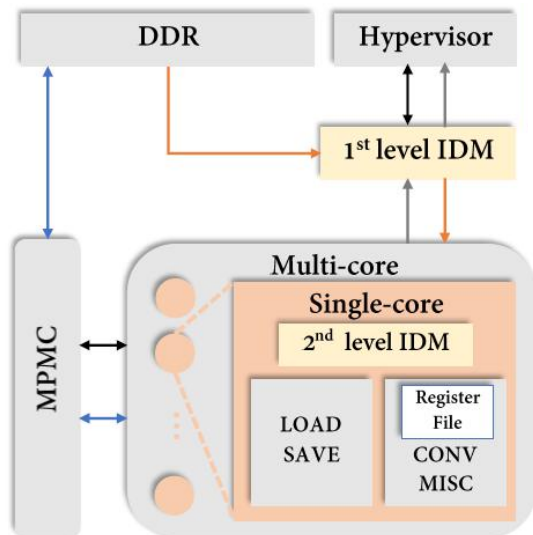
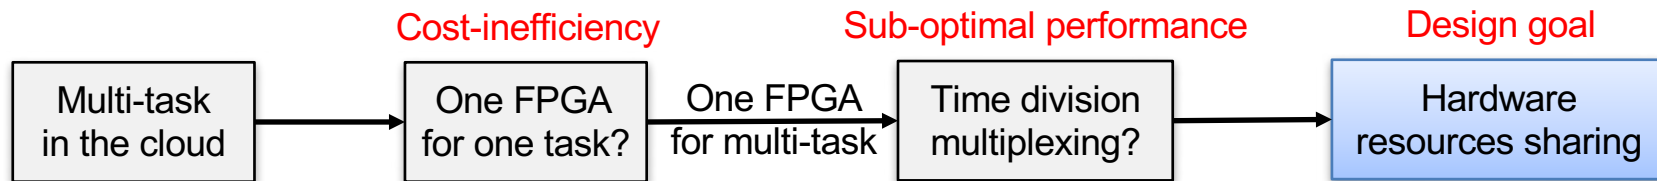
- Time sharing:
  - Hard to ensure physical isolation
  - Sub-optimal performance
- Space sharing:
  - Easy to ensure isolation
  - High reconfiguration overhead

# How to Support Multiple Tasks in the Cloud?



Normalized performance with different hardware resources in the single-task and static-workload scenario.

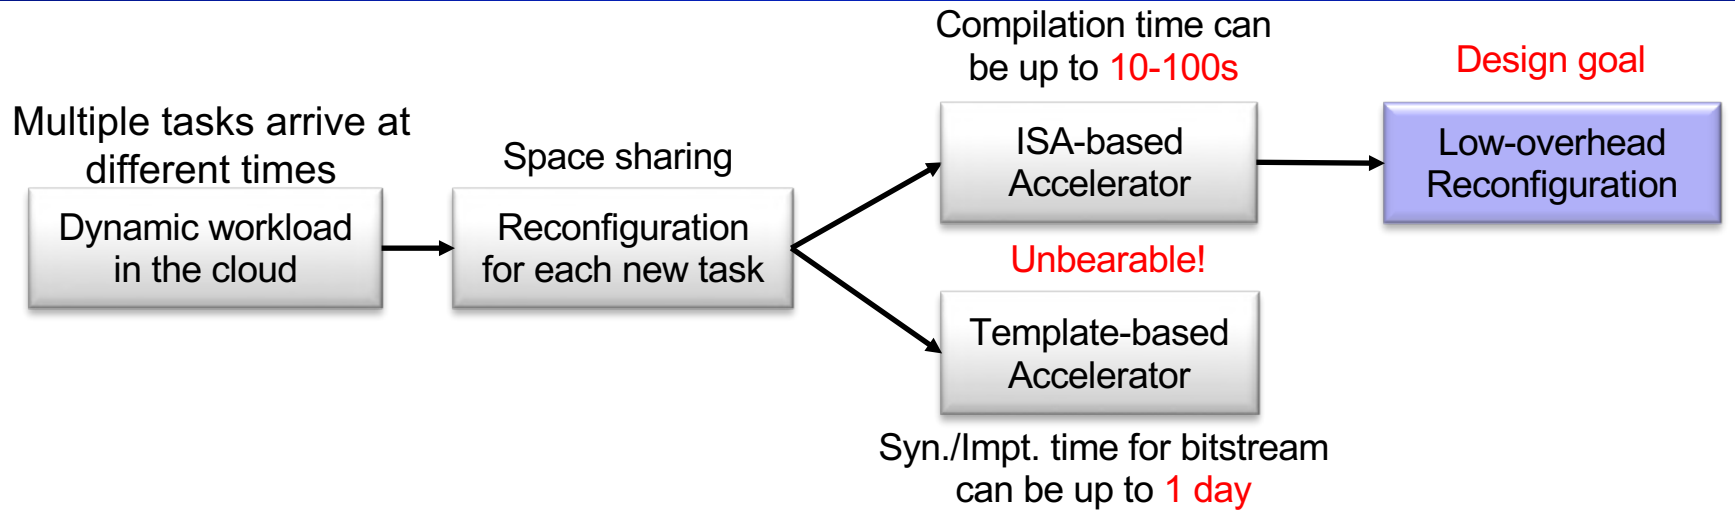
# How to Support Multiple Tasks in the Cloud?



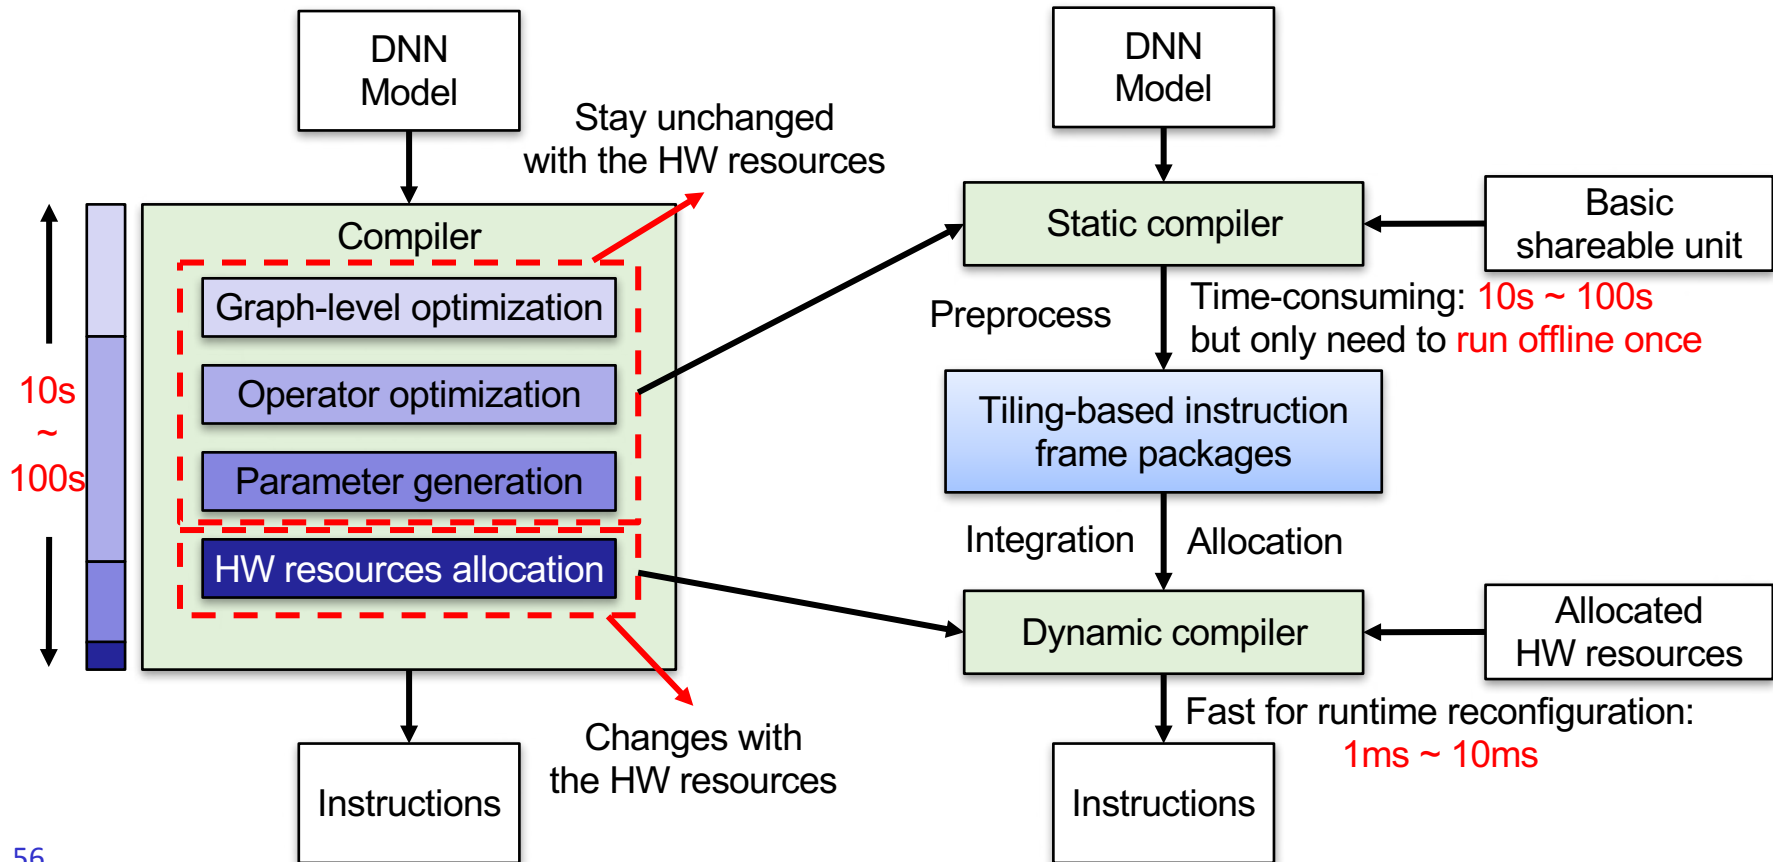
Our solution: Multi-core CNN Accelerator

- Good isolation: each core is an unique hardware resource
- Dynamic reconfiguration: allocate the computing cores based on the CNN workload
- Optimized for both latency and throughput
- Key problem: **high reconfiguration overhead?**

# How to Support Dynamic Workload in the Cloud?



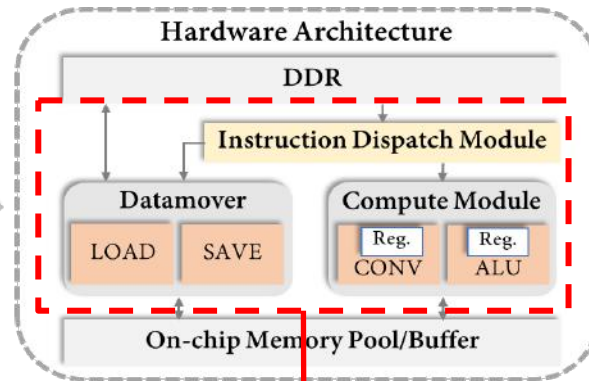
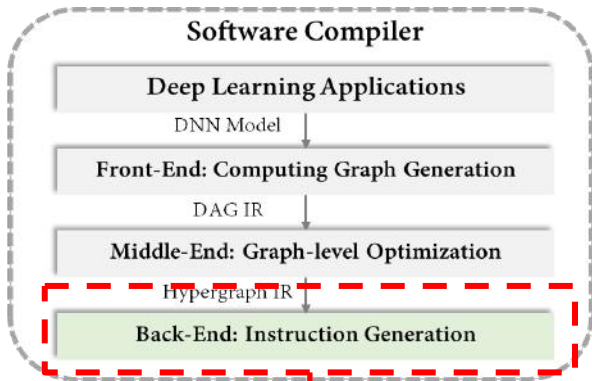
# Low-overhead Reconfiguration of ISA-based Accelerator






# Design Techniques

Baseline design  
of ISA-based  
CNN accelerator



Design goal for software:  
**Low-overhead reconfiguration**



Design goal for software:  
**Multi-Task concurrent execution**



Design Techniques  
for Virtualization

Instruction Support:  
A **bridge** between the static  
and dynamic compiler.

Instruction Support:  
Original IDM: Module-level scheduler  
New IDM: **Task-level** scheduler



# Experiments

- Experiment Setup

- Hardware platform: Xilinx VU9P FPGA (Vivado 2018.2) + Intel Xeon E5-2643 CPU
- Deep learning model (Caffe): VGG-16, ResNet-18, MobileNet-v1
- **Simulation of applications in the cloud: Poisson distribution [1]**
- CNN accelerator: Parallelism is 8192 ops/cycle (2048 DSPs, each for 2xINT8 MAC), running at 200MHz.
- Static implementation: 1) single large core (1 x 8192). 2) eight small cores (8 x 1024)
- Virtualized implementation: Multi-core virtualized design (16 x 512).

- Hardware resources utilization

- The virtualized multi-core design utilize **3.25%** more BRAMs, **3.08%** more LUTs than the static multi-core design.
- Noted that the resources of URAM are not utilized in this experiment.

Implementation	LUT	FF	BRAM	DSP
<b>Static</b> single core	22.01%	11.06%	35.23%	29.94%
<b>Static</b> multi-core	54.51%	13.3%	86.34%	29.94%
<b>Virtualized</b> multi-core	57.59%	30.59%	89.59%	29.94%

[1] Hamzeh Khazaei et al. 2012. Performance of cloud centers with high degree of virtualization under batch task arrivals. IEEE Transactions on Parallel and Distributed Systems.





# Experiments

Evaluation on the *isolation characteristics* of the **Public Cloud**:

- Effective isolation: the average performance difference is **<1 ms**.

Evaluation on the *single-task performance (latency)* of the **Private Cloud**:

- The average performance loss of multi-core is **less than 3%**.



# Experiments

Evaluation on the reconfiguration overhead:

- Complete re-generation of instructions can be **12-200s**, which is unbearable.
- Time sharing method has no context-switch cost, but it is cost-efficient and sub-optimal.
- Our space sharing method introduces **~1ms** context-switch overhead, but achieves better performance on throughput and negligible performance loss on latency.

**10.6-35.2%** extra overhead of static compilation, which only needs to **run offline once**.


CNN Model	Original Compilation	Static Compilation	Dynamic Compilation	Transfer Time	Context-switch Cost	Computation Time	
						Min(8192)	Max(512)
VGG-16	201930	223457	1.75	0.26	2.04	21.2	219.1
ResNet-18	25336	34257	1.55	0.09	1.68	8.4	37.7
MobileNet-v1	12074	15767	0.81	0.09	0.94	6.0	15.2

**~1ms** context-switch cost is **acceptable** for the computation time of **6.0-219.2ms** during runtime.

# Experiments

Evaluation on the performance of *throughput* on the *multi-task* scenario:

- Static single core(8192): low throughput for multiple tasks due to only one core.
- Static multi-core(8x1024): the worst throughput when there is only one task.
- Virtualized multi-core (16x512): the best throughput for different workloads.



1.88x ~ 2.20x  
higher throughput

## Future Work

---

- Implementation on the cloud FPGA (Alibaba Cloud)
- Further improvement on the dynamic compilation cost
  - Down to 0.1ms, 1% of the inference time
- Further improvement on the hardware architecture.



## Case Study 3: Reliable NN on FPGA?

# When NN Chips Meet Fault Tolerant Problem

## Environment Effect



Outer Space  
Radiation



Temperature  
Variation




Ultra Low  
Power Supply



## Inherent Property



Device Aging



Immature New  
Device Tech.



## Artificial



Hacker  
Attack



Reliability Problem

Security Problem

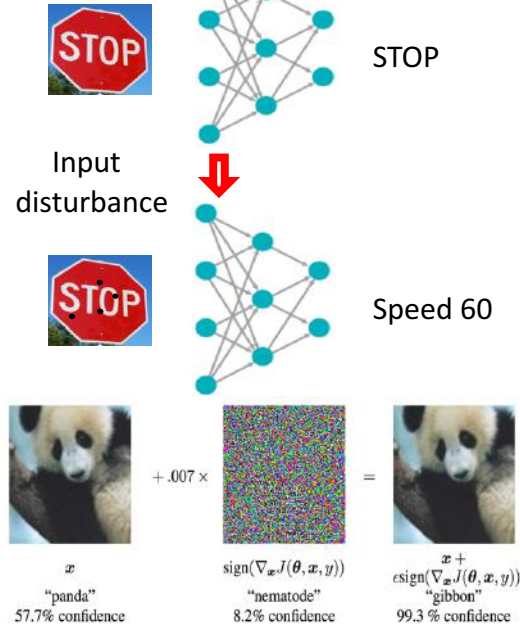




# Reliability problems of NN system

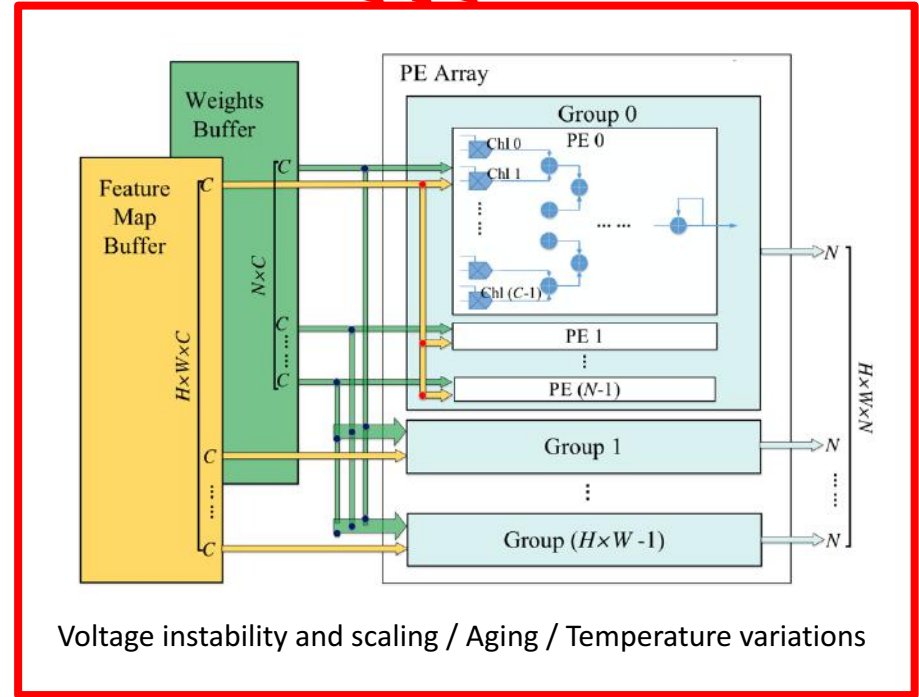
## Algorithm:

Adversarial Robustness, Neural networks are susceptible to subtle input disturbances and output incorrect results



Hardware-related:

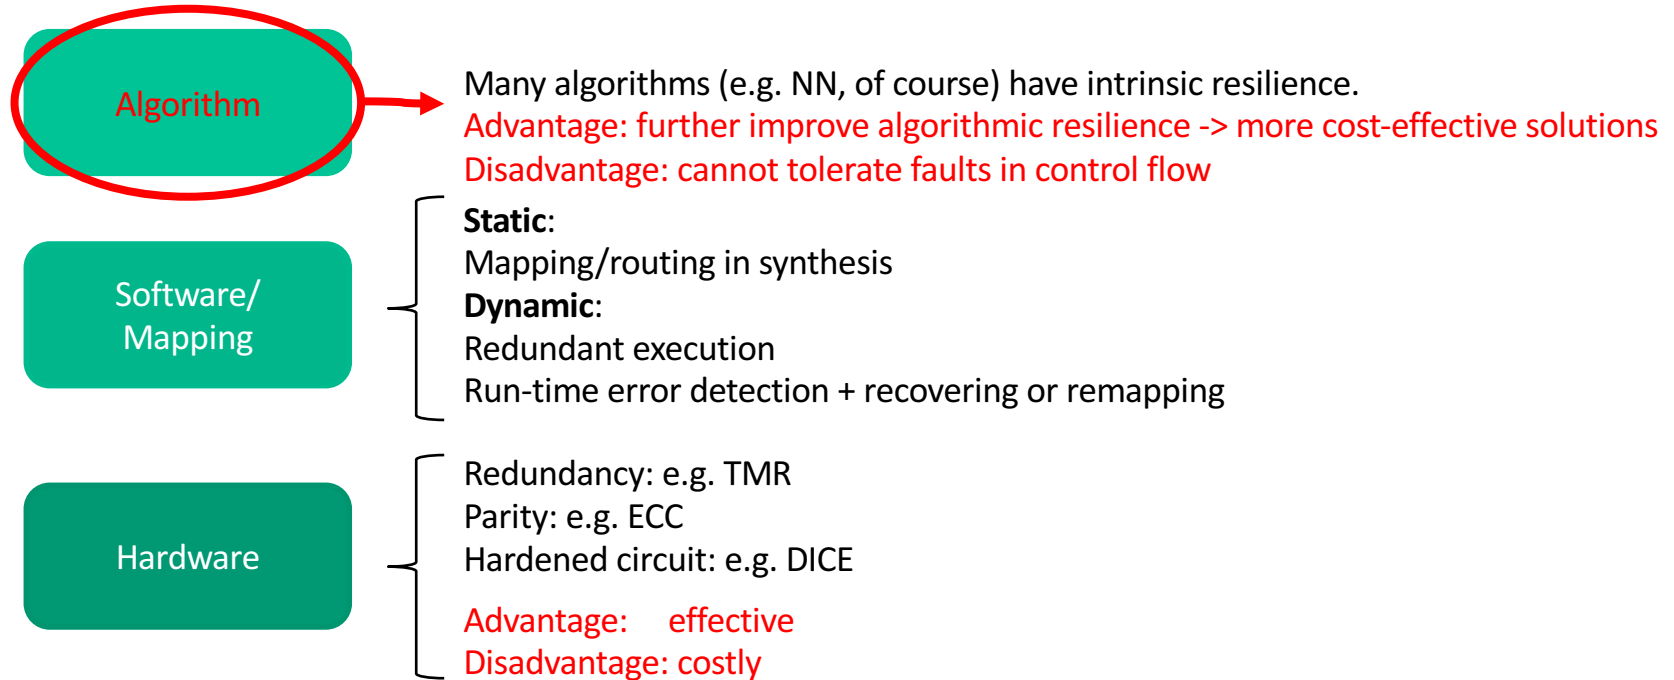
Atmospheric neutrons  
Radioactive impurities



Voltage instability and scaling / Aging / Temperature variations

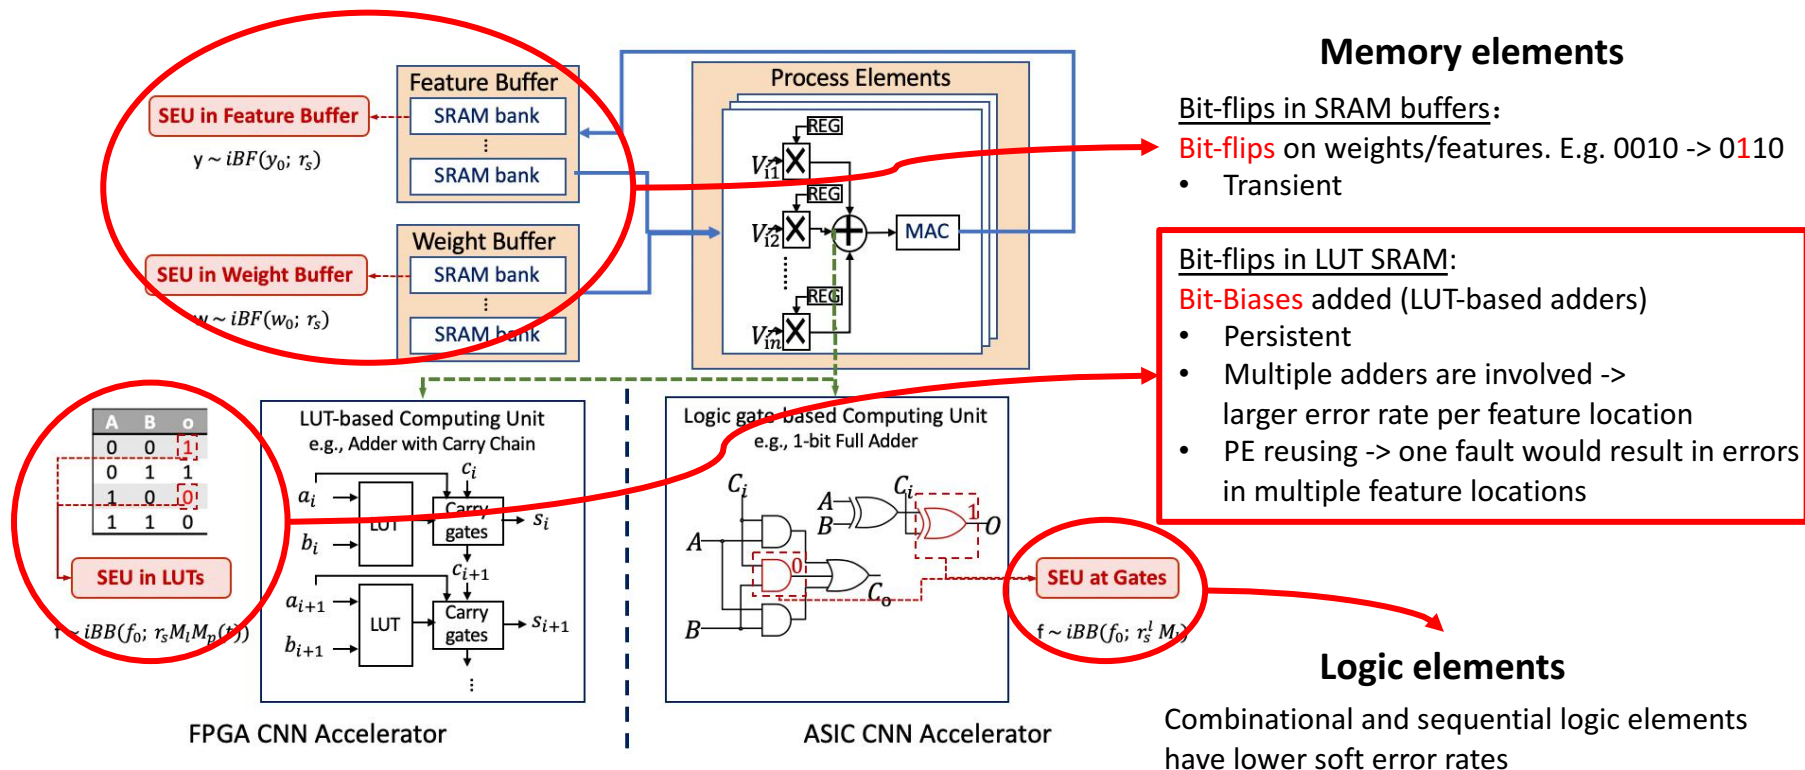
Computational Faults

# Mitigations for computational faults



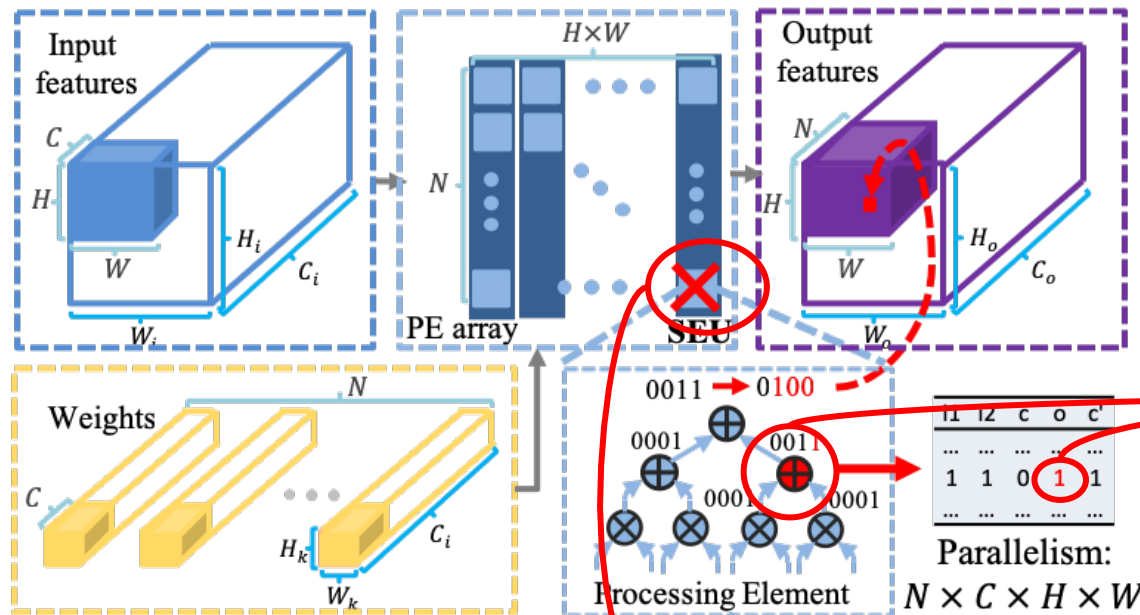
Most studies on fault-tolerance are **exploiting** the NN's algorithmic resilience.  
Few studies are **improving** the NN's algorithmic resilience.

# Analysis of Platforms: FPGA and ASIC



# (LUTs) Feature Fault Model: MAC-i.i.d random Bit-Bias

## MiBB feature fault model



**Random Bit-Bias Faults:** Error added as biases

$$y = g(W \circledast x + b + \theta \cdot 2^{\alpha-l} \cdot (-1)^\beta)$$

$$\text{s.t. } \theta \sim \text{Bernoulli}(p)^{C_o \times H \times W}$$

$$\alpha \sim U\{0, \dots, Q-1\}^{C_o \times H \times W}$$

$$\beta \sim U\{0, 1\}^{C_o \times H \times W}$$

$$p = p_m \times ck^2$$

$p_m$  is the prob that an adder encounter error

$\alpha$  refers to the fault bit

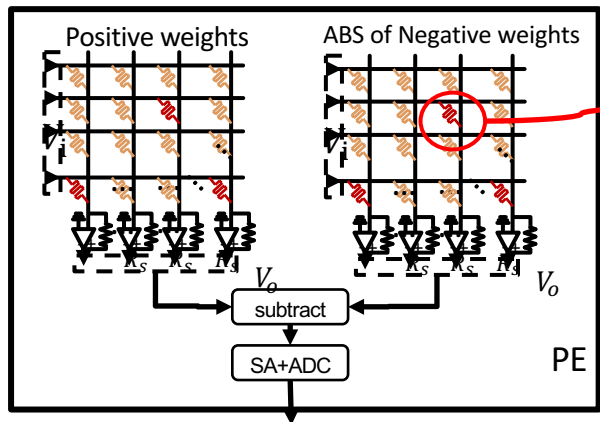
$\beta$  refers to the sign of biases

$ck^2 - 1$  additions spatially expanded onto adders

$\theta_i$  refers to whether a fault occurred in one feature value

# (RRAM) Weight Fault Model: arbitrary-distributed SAF

## RRAM with multi-bit cells



Weight representation range:  $[-R^w, R^w] = [-2^{-l}(2^{Q+1} - 1), 2^{-l}(2^{Q+1} - 1)]$

stuck at 0 with  $p_0$   
stuck at Maximum  $R^w$  with  $p_1$

## Q-bit adSAF weight fault model

Computation of one layer:


$$y = g(W' \otimes x + b)$$

$$\text{s.t. } W' = (1 - \theta) \cdot W + \theta \cdot e$$

$$\begin{cases} \theta \sim \text{Bernoulli}(p_0 + p_1)^{C_o \times c \times k \times k} \\ m \sim \text{Bernoulli}\left(\frac{p_1}{p_0 + p_1}\right)^{C_o \times c \times k \times k} \\ e = R^w \text{sgn}(W) \cdot m \end{cases}$$

$\theta$  refers to the weight fault position mask  
 $m$  refers to the fault type mask, SAF0/SAF1  
 $e$  refers to the fault target value,  $0/\pm R^w$

## RRAM with single-bit cells (e.g. 8-bit weights)



memristor stuck at 1/0: a bit of a weight stuck at 0/1

$\theta$  refers to the bit fault position mask  
 $e$  refers to the target value at each bits, 0/1

## 1-bit adSAF weight fault model

Change of one weight:

$$w' = \text{sgn}(w) 2^{-l} (((-\theta) \wedge 2^l |w|) \vee (\theta \wedge e))$$

$$\theta = \sum_{q=1}^Q \theta_q 2^{q-1}$$

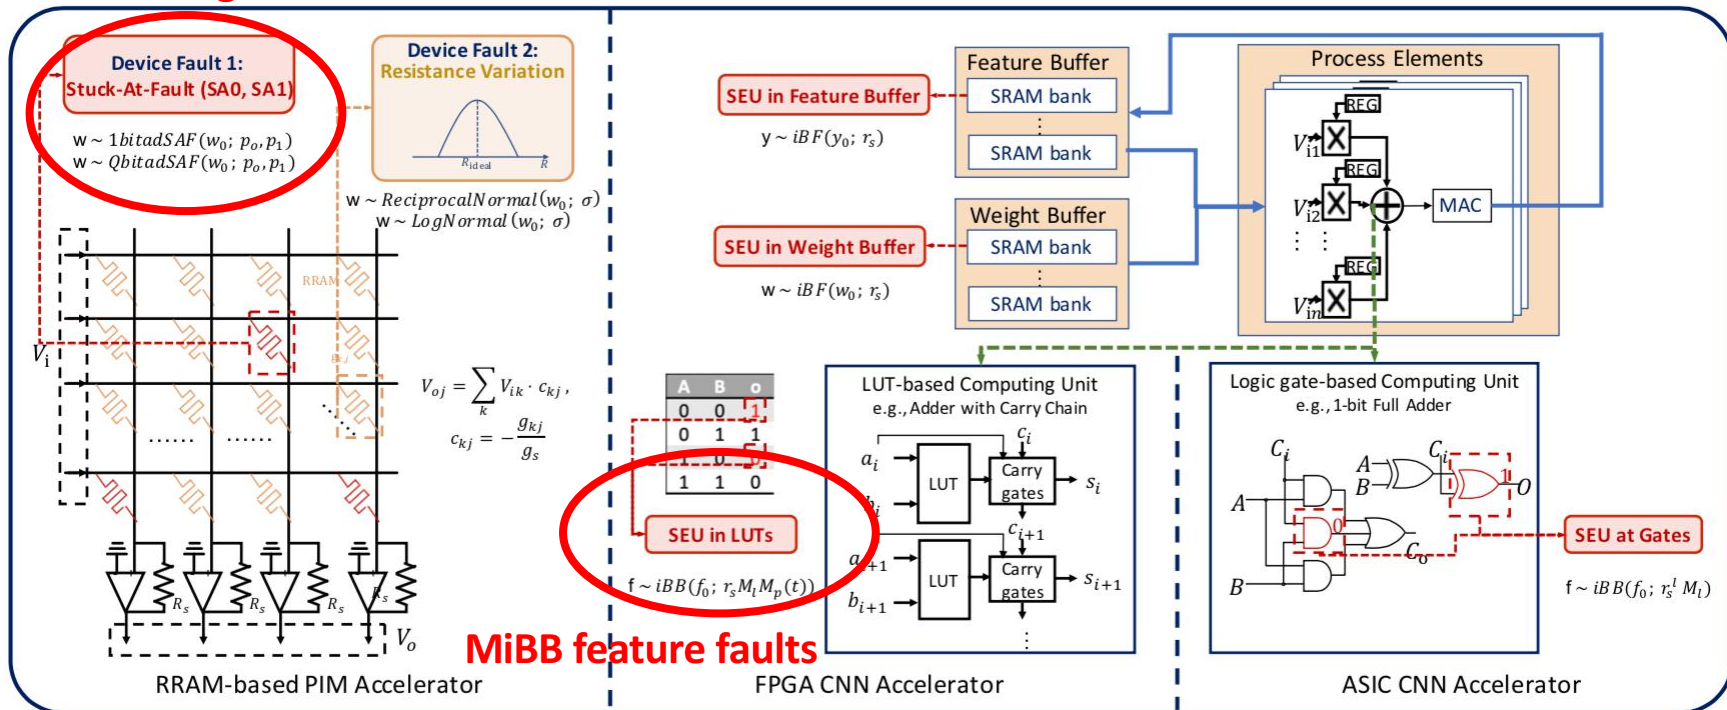
$$\theta_q \stackrel{iid}{\sim} \text{Bernoulli}(p_0 + p_1), \quad q = 1, \dots, Q$$

$$e = \sum_{q=1}^Q m_q 2^{q-1}$$

$$m_q \stackrel{iid}{\sim} \text{Bernoulli}\left(\frac{p_1}{p_0 + p_1}\right), \quad q = 1, \dots, Q$$

# Summary of Computational Faults

## adSAF weight faults



# Influence of Neural Architectures on Reliability

Preliminary experiments show that **different archs have different resilience capabilities**.

**Feature faults: Random Bit-Bias**

Comparison with different  $p_m$  ( $0, 10^{-5}, 10^{-4}$ )

Model	Acc( $0/10^{-5}/10^{-4}$ )	#Params	#FLOPS
ResNet-20	94.7/63.4/10.0	11.2M	1110M
VGG-16 <sup>†</sup>	93.1/21.4/10.0	14.7M	626M
MobileNet-V2	92.3/10.0/10.0	2.3M	182M

<sup>†</sup>: For simplicity, we only keep one fully-connected layer of VGG-16.

**Weight faults: 8-bit SAF**


Comparison with different SAF ratios ( $0, 4\%, 8\%$ )

Model	Acc( $0/4\%/8\%$ )	#Params	#FLOPS
ResNet-20	94.7/64.8/17.8	11.2M	1110M
VGG-16	93.1/45.7/14.3	14.7M	626M
MobileNet-V2	92.3/26.2/11.7	2.3M	182M

Operation types and connections are possible to affect the reliability of models. For example:


MobileNet  
Block

×

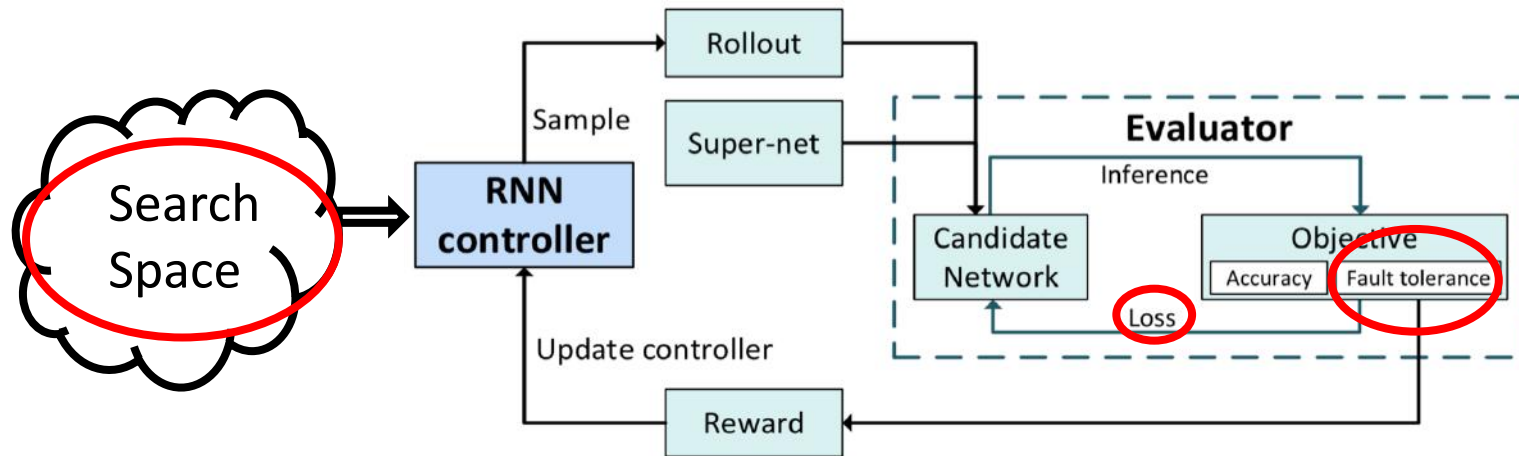


ResNet  
Block

✓



# NAS Problem formalization



$$\min_{\alpha \in \mathcal{A}} E_{x_v \sim D_v} [E_{f \sim F} [R(x_v, \text{Net}(\alpha, w^*(\alpha)), f)]]$$

$$\text{s.t. } w^*(\alpha) = \operatorname{argmin}_w E_{x_t \sim D_t} [E_{f \sim F} [L(x_t, \text{Net}(\alpha, w), f)]]$$

$F$ : Distribution characterized by the fault model

$D_t$ : Training dataset

$D_v$ : Valid dataset

$\mathcal{A}$ : Architecture space

$\alpha$ : An architecture rollout sampled from the architecture space

$w$ : network weights

$\text{Net}(\alpha, w)$ : Candidate Network

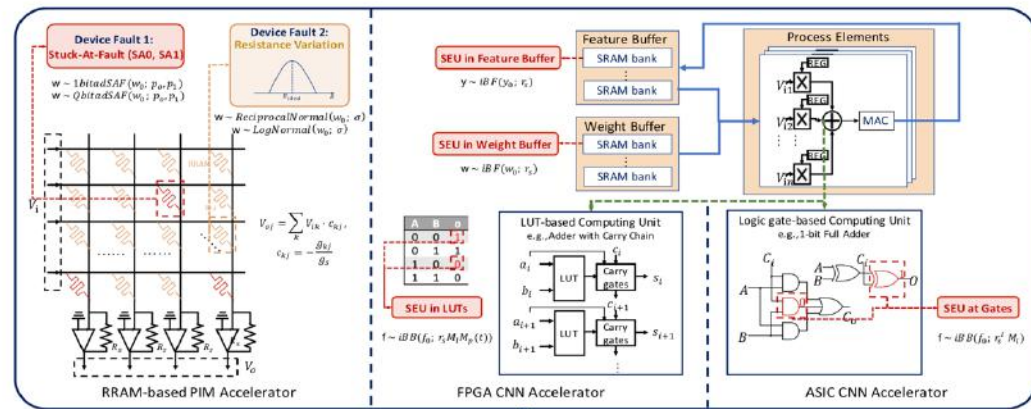
$w^*(\alpha)$ : optimal network weights of  $\alpha$  (get on the training dataset)

$L$ : Loss

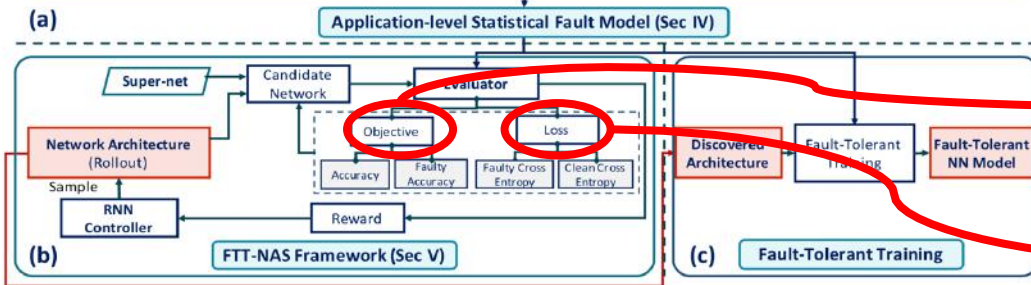
$R$ : Reward



# The Overall Workflow



- Analyze and formalize the statistical fault models  
 -> fault models (guide fault injection)
- Run neural architecture search  
 -> architecture
- Fault tolerant training  
 -> fault-tolerant model



reward

$$R = (1 - \alpha_r) * \text{acc}_c + \alpha_r * \text{acc}_f$$

loss

$$L = (1 - \alpha_l) * \text{CE}_c + \alpha_l * \text{CE}_f$$

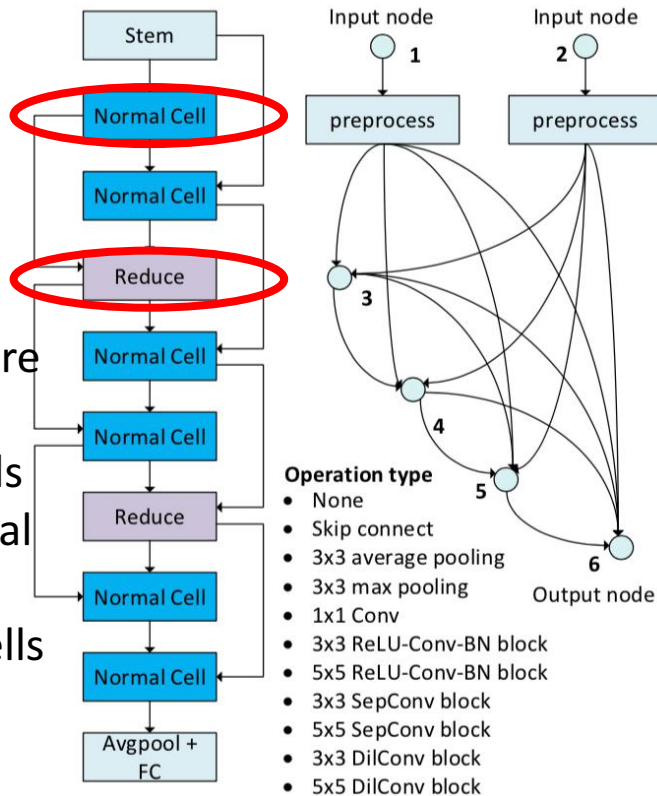
FTT-NAS:  $\alpha_l > 0, \alpha_r > 0$ . search with both fault-tolerant reward and fault-tolerant loss

FT-NAS (a degraded version):  $\alpha_l = 0, \alpha_r > 0$  search with fault-tolerant reward and without fault-tolerant loss

# Search Space Design


Artificial  
Macro  
Architecture

Divide Cells  
into Normal  
Cells and  
Reduce Cells



Each cell has four nodes

Every node selects two  
operations from all its  
predecessor nodes






# Results: adSAF weight faults

## Primitive Preference:


Normal convolutions > separable/dilation convolutions

## Inspection of weight distribution of different operations




## Hypothesis:

Larger weight magnitudes -> Larger deviation caused by SAFs



(a)



(b)

TABLE V: Comparison of different architectures under the adSAF weight fault model

Arch	Training	Clean accuracy	Accuracy with weight faults (%)					#FLOPS	#Params
			0.04	0.06	0.08	0.10	0.12		
ResNet-20	clean	<b>94.7</b>	<b>64.8</b>	<b>34.9</b>	17.8	12.4	11.0	1110M	11.16M
VGG-16	clean	93.1	45.7	21.7	14.3	12.6	10.6	626M	14.65M
MobileNet-V2	clean	92.3	26.2	14.3	11.7	10.3	10.5	182M	2.30M
<b>W-FTT-Net-20</b>	clean	91.7	54.2	30.7	<b>19.6</b>	<b>15.5</b>	<b>11.9</b>	1020M	3.05M
ResNet-20	$p=0.08$	92.0	86.4	77.9	60.8	41.6	25.6	1110M	11.16M
VGG-16	$p=0.08$	91.1	82.6	73.3	58.5	41.7	28.1	626M	14.65M
MobileNet-V2	$p=0.08$	86.3	76.6	55.9	35.7	18.7	15.1	182M	2.30M
<b>W-FTT-Net-20<sup>†</sup></b>	$p=0.08$	90.8	86.2	79.5	69.6	53.5	38.4	919M	2.71M
<b>W-FTT-Net-40</b>	$p=0.08$	<b>92.1</b>	<b>88.8</b>	<b>85.5</b>	<b>79.3</b>	<b>69.2</b>	<b>54.2</b>	3655M	10.78M

<sup>†</sup>: The “- $N$ ” suffix means that the base of the channel number is  $N$ .

# Results: adSAF weight faults

NN model trained under 8bit-adSAF model, can also tolerate 1bit-adSAF faults, i.i.d random Bit-Flip faults




Fig. 9: Accuracy curve under different weight fault models. (a)W-FTT-Net under 8bit-adSAF model (b)W-FTT-Net under 1bit-adSAF model. (c)W-FTT-Net under iBF model

# Future Work

---

- Fault models
  - Now the modeling corresponds to **template-based FPGA accelerators**, for **instruction-based FPGA accelerators**, the fault model would be different since other factors should be considered:
    - Unrolling in kernel's spatial/channel dimensions leads to temporal reusing of hardware components, thus would leads to different fault propagation behavior for one feature value
    - Unrolling in OFM (output feature map) spatial/channel dimensions leads to repetitive error pattern in the OFM




## Case Study 4: Secure NN on NVM/FPGA?



# Analysis for NN Fault Problems



What's the general characteristics for various NN



Model	Dataset	Classification Accuracy				Encryption Config.
		Encrypted		Baseline		
		Top-1	Top-5	Top-1	Top-5	
ResNet-18	ImageNet	0.704% (-69.05%)	2.452% (-86.62%)	69.75%	89.07%	$N=20, \epsilon=0.2$
ResNet-50	ImageNet	0.438% (-75.69%)	1.540% (-91.32%)	76.13%	92.86%	$N=30, \epsilon=0.1$
ResNet-101	ImageNet	0.144% (-77.24%)	0.758% (-92.78%)	77.38%	93.54%	$N=20, \epsilon=0.1$
VGG-16	ImageNet	0.818% (-70.77%)	3.478% (-86.90%)	71.59%	90.38%	$N=30, \epsilon=0.2$

If we disturb  $N=20\sim30$  picked parameters ( $< 0.1\%$ ) each layer, NN will failed its function.


**KEY POINT:** We observed that, most NN networks have a fraction of nodes (**key point**) which influence the performance a lot.



# The Confidentiality Protection of NVM-based NN Systems


## Memory Wall

Bottleneck of speed and energy




## Compute-in-Memory

Merge the mem. & compute




## Non-volatility

Weights will persist in memory



Key Point  
Encryption



## Heavy Weights

Cause large encryption overhead

01001...10010010100...  
↓ Encryption  
11011...10010101001...

## Motivating Encryption

To keep the data secure


## Confidentiality Attacks

Deployed NN models may be leaked

# How We Distinguish the Key Points of NN Weights

## Encryption Goal

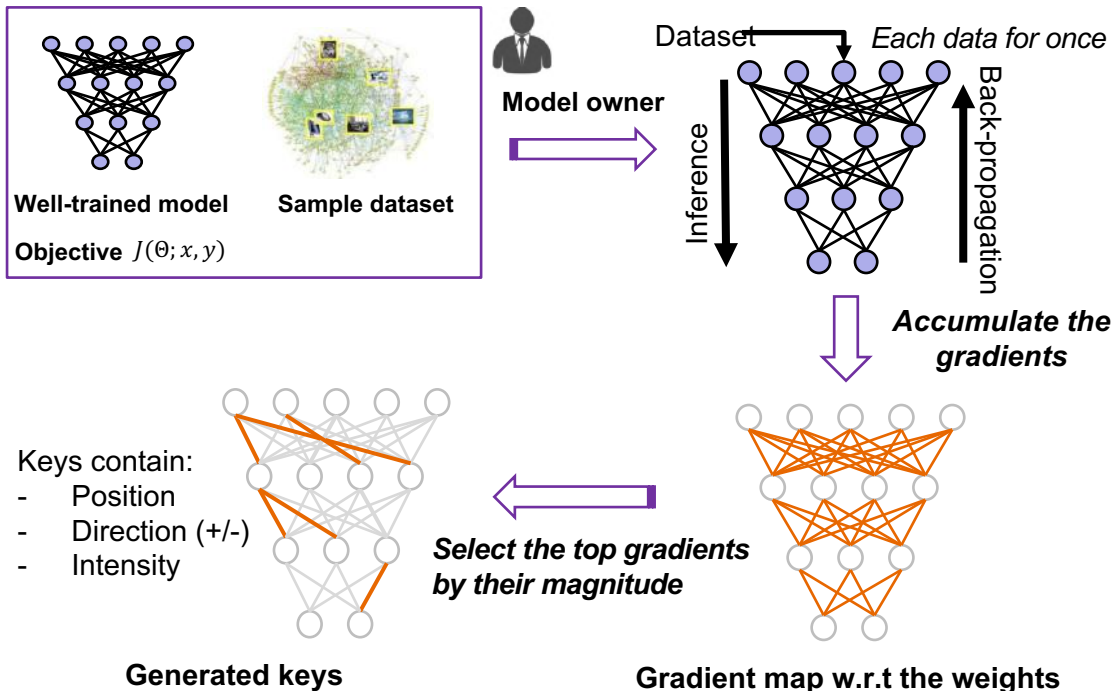
Encrypting key weights  
to disable the NN




$$\nabla_{\Theta} f = \lim_{\varepsilon \rightarrow 0} \frac{f(\Theta + \varepsilon) - f(\Theta)}{\varepsilon}$$

**Gradient** shows the speed of  
objective rising


## Our Proposal Sparse Fast Gradient Encryption (SFGE)



# Experimental Results: the Effectiveness and Efficiency



Accuracy versus # encryption  $N$

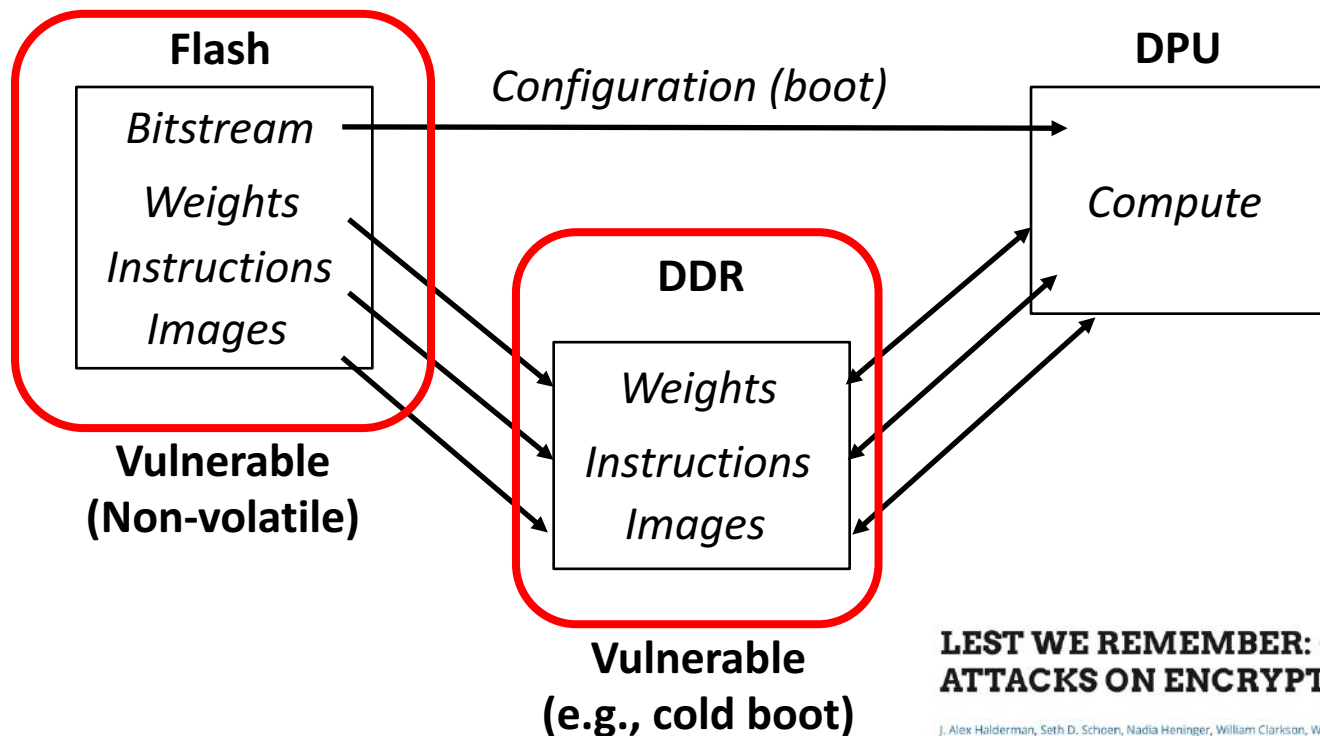


VGG-16 accuracy versus # encryption  $N$   
under different intensity  $\epsilon$

**Only Encrypting 20~30 Weights per Layer in the Neural Network can Strictly Protect the model Confidentiality**

# Does FPGA design secure enough?

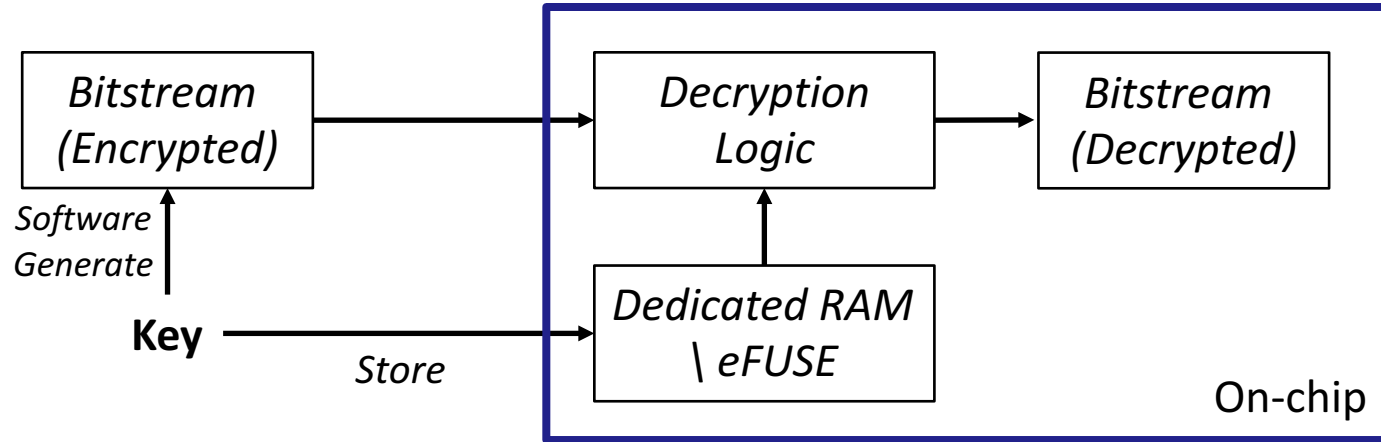
## A typical workflow of FPGA neural network accelerator



**LEST WE REMEMBER: COLD BOOT  
ATTACKS ON ENCRYPTION KEYS**

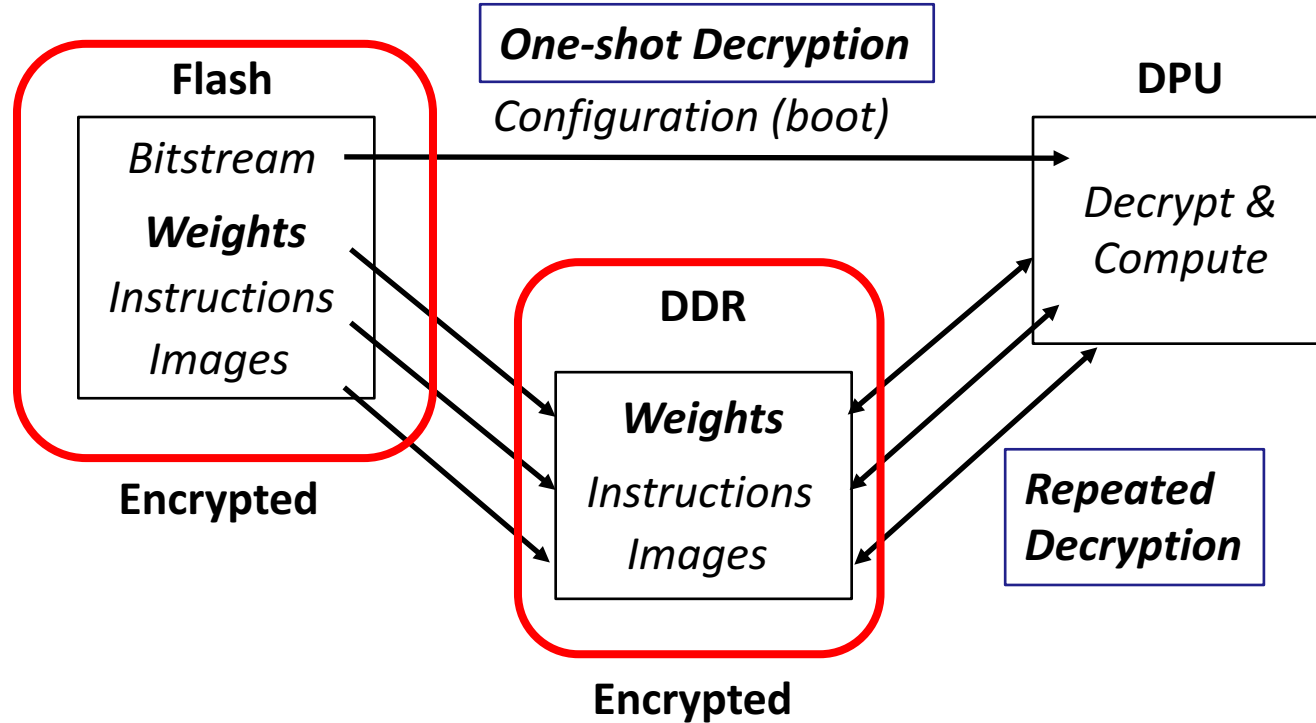
J. Alex Halderman, Seth D. Schoen, Nadia Heninger, William Clarkson, William Paul, Joseph A. Calandrino, Ariel J. Feldman, Jacob Appelbaum, and Edward W. Felten.

# Current process to protect security of FPGA bitstreams



However, how to secure **the deployed NN models?**

# Our method can be transferred to secure FPGA NN designs



**Decrypting part of the weights – lower latency, more secure**

# Summary

---

## **What we have or we can expect (power efficiency, computation power)**

- AI platforms are covering everywhere.
- We already have many methods to make the chips powerful.


## **What's next: (user friendly, customizable, low cost, real-time, reliable, secure...)**

- Are the chips ready for real applications?
- Can we use the large 'TOPs' to the extreme?
- Are the systems robust enough?
- Are the systems secure enough?

**Learning capability: application, algorithm, and hardware arch**

# Users care much more than energy efficiency

Flexibility, Security, Robustness, ...





# Acknowledgement

**To my students and collaborators**



**To supporting companies and agencies**





# Thanks for your attention!

yu-wang@tsinghua.edu.cn

